

SA2 - Implementing YOLOv26 with NMS-Free Detection

Barce, John Bernard B.

Respecio, Luis Arnold N.

Overview

In this section, the environment setup for the model training process is detailed. The first step involves mounting Google Drive to the Colab runtime, enabling access to the necessary files stored in the user's Google Drive. It also includes the installation of required packages, such as ultralytics, which are essential for the YOLOv26 model training. The environment setup imports several Python libraries: os for file and directory operations, shutil for high-level file management tasks, yaml for working with configuration files, numpy for numerical operations, and ultralytics for model-specific checks. The ultralytics.checks() function is executed to ensure that all required components for the training are properly configured. This setup ensures that the runtime environment is correctly prepared to run the object detection model (Google Colab, 2023).

Section 1 - Environment Setup

This section details the initialization and training setup for three different models. First, a pandas DataFrame is created to organize the hyperparameter configurations for all three models, allowing for easy comparison. Key hyperparameters include the model name, input image size (set to 640x640 for a balance of speed and accuracy), optimizer (AdamW, SGD, and auto), batch size (with automatic selection for certain configurations), and learning rate (with varying values to test sensitivity). The table displaying these configurations is formatted for easy readability and includes a footer separator for visual clarity. This section ensures that all training configurations are well-documented and easily accessible before proceeding with the model training process (Pandas Documentation, 2023).

```
# Import the Google Colab drive module to enable access to Google Drive
# This module provides functionality for mounting cloud storage
from google.colab import drive

# Mount Google Drive to the Colab runtime at the specified path!pip
install ultralytics --quiet
# This allows the notebook to read and write files from Google Drive#
The --quiet flag suppresses detailed installation output

drive.mount('/content/drive')
```

```
Drive already mounted at /content/drive; to attempt to forcibly
remount, call drive.mount("/content/drive", force_remount=True).

# Import os module for operating system dependent functionality
# Provides functions for file and directory operations
import os

# Import shutil module for high-level file operations
# Used for copying and removing files and directory trees
import shutil

# Import yaml module for reading and writing YAML configuration files
# Enables parsing and dumping of YAML formatted data
import yaml

# Import numpy for numerical operations and array manipulation
# Provides support for large multi-dimensional arrays and matrices
import numpy as np

ultralitics.checks()

# Import pandas for data manipulation and analysis# Validates that all
required components are properly configured
# Provides DataFrame structures for organizing tabular data# Run
ultralitics system checks to verify installation and dependencies

import pandas as pd

print(f"Ultralytics version: {ultralitics.__version__}")

# Import matplotlib.pyplot for creating visualizations and plots#
Helps verify the correct version is installed for compatibility
# Provides MATLAB-like plotting interface for generating charts# Print
the current version of the ultralytics library being used

import matplotlib.pyplot as plt

warnings.filterwarnings("ignore")

# Import matplotlib.image for reading and displaying image files#
Suppresses warning messages to keep output clean
# Enables loading images into numpy arrays for visualization#
Configure warnings to be ignored and not displayed

import matplotlib.image as mpimg

import torch, gc

# Import YOLO class from ultralytics for object detection# torch
provides neural network functionality, gc manages memory cleanup
# Provides the main interface for loading and using YOLO models#
```

```

Import torch for PyTorch deep learning framework and gc for garbage
collection

from ultralytics import YOLO

import warnings

# Import ultralytics package to access version information# Allows
filtering and suppression of warning messages
# Contains the complete YOLO implementation and utilities# Import
warnings module to control warning message display

import ultralytics

Ultralytics 8.4.18 □ Python-3.12.12 torch-2.10.0+cu128 CUDA:0 (Tesla
T4, 14913MiB)
Setup complete □ (2 CPUs, 12.7 GB RAM, 44.0/112.6 GB disk)

# Import shutil and os modules for file system operations
# These modules enable copying directories and checking paths
import shutil, os

# Define the source path pointing to the dataset in Google Drive
# This is where the original dataset files are stored
SRC = "/content/drive/MyDrive/DataScience 3"

# Define the destination path for the local copy of the dataset
# This is where files will be copied for faster access during training
DST = "/content/dataset"

# Check if the destination directory already exists from a previous
run
# This prevents issues with incomplete or outdated copies
if os.path.exists(DST):
    # Remove the entire directory tree at the destination path
    # This ensures a clean slate before copying new data
    shutil.rmtree(DST)

    # Print confirmation that the old directory was removed
    # Provides feedback to the user about the cleanup operation
    print("Removed old incomplete copy.")

# Print status message indicating the copy operation is starting
# Informs the user that the dataset transfer is in progress
# Shows the organization and size of each dataset split
print("Copying dataset to Colab local storage...")

# This transfers all dataset files to local storage for faster access
# This provides information about dataset size
shutil.copytree(SRC, DST)

```

```

# Count the number of files in the subdirectory
# Ensures we only count files in actual directories
# Confirms that all files have been transferred
for item in sorted(os.listdir(DST)):
    item_path = os.path.join(DST, item)    # Create the full path for
    the subdirectory

    # Iterate through each item in the destination directory in sorted
    order
    # Enables checking its properties and contents
    if os.path.isdir(item_path):    # Check if the current item is a
    directory
        # Only process directories to explore their contents
        # This examines the structure of the copied dataset
        for sub in sorted(os.listdir(item_path)):    # Iterate through
        subdirectories within the current directory
            sub_path = os.path.join(item_path, sub)    # Create the
            full path by joining the destination with the item name

            if os.path.isdir(sub_path):    # Check if the subdirectory
            is actually a directory
                pass    # This ensures we only count files in actual
                directories

# Print completion message to indicate successful copy operation
# Confirms that all files have been transferred
print("Done!")

```

```

Removed old incomplete copy.
Copying dataset to Colab local storage...
Done!
runs/yolo26n_E25_AdamW_B4_LR001/ (7 files)
test/images/ (247 files)
test/labels/ (247 files)
train/images/ (1676 files)
train/labels/ (1676 files)
valid/images/ (465 files)
valid/labels/ (465 files)

```

```

# Define the base path where the dataset is located
# This path will be used to construct paths to train, validation, and
test sets
DATASET_PATH = "/content/dataset"

```

```

# Create a dictionary containing dataset configuration for YOLO
# This structure defines paths and class information needed for
training
data_yaml = {
    # Define the path to training images directory
    # YOLO will use these images to train the model

```

```

"train": os.path.join(DATASET_PATH, "train", "images"),

# Define the path to validation images directory
# YOLO will use these images to validate during training
"val": os.path.join(DATASET_PATH, "valid", "images"),

# Define the path to test images directory
# YOLO will use these images for final evaluation
"test": os.path.join(DATASET_PATH, "test", "images"),

# Specify the number of classes in the dataset
# This dataset contains 4 different object classes
"nc": 4,

# List the names of all classes in the dataset
# These labels correspond to the objects the model will detect
"names": ["crossing", "jaywalking", "pedestrian", "waiting"],
}

# Define the full path for the YAML configuration file
# This file will be saved in the dataset directory
DATA_YAML = os.path.join(DATASET_PATH, "data_local.yaml")

# Open the YAML file in write mode
# This creates or overwrites the configuration file
with open(DATA_YAML, "w") as f:
    # Write the data_yaml dictionary to the file in YAML format
    # default_flow_style=False ensures readable multi-line format
    yaml.dump(data_yaml, f, default_flow_style=False)

# Print a header indicating the YAML file contents will be displayed
# The checkmark provides visual confirmation of success
print("✅ data_local.yaml:\n")

# Open the YAML file in read mode
# This allows reading back the contents to verify
with open(DATA_YAML, "r") as f:
    # Read and print the entire contents of the YAML file
    # This confirms the configuration was written correctly
    print(f.read())

# Iterate through each dataset split to verify image counts
# This validates that all required dataset directories exist
for split in ["train", "val", "test"]:
    img_path = data_yaml.get(split, "") # Get the image path for the
    current split from the dictionary
    # Check if the image directory exists on the filesystem
    if os.path.exists(img_path):
        # Count the number of image files in the directory
        count = len([f for f in os.listdir(img_path) if

```

```
f.lower().endswith((".jpg", ".jpeg", ".png"))])
    # Print the split name and count with a success checkmark
    print(f" {split}: {count} images")
else:
    # Print an error message if the directory doesn't exist
    print(f" {split}: NOT FOUND")

data_local.yaml:

names:
- crossing
- jaywalking
- pedestrian
- waiting
nc: 4
test: /content/dataset/test/images
train: /content/dataset/train/images
val: /content/dataset/valid/images

train: 1676 images
val: 465 images
test: 247 images
```

Section 2 - Model Initialization & Training

This section details the process of evaluating the trained models through the use of confusion matrices. First, the path to the dataset configuration file is defined, allowing the models to access the necessary data. Each model is associated with a specific output directory, where the training results, including the best model weights, are stored. The section then verifies if the best model weights file (best.pt) exists in the respective directories for each model. It iterates through each model's directory, checking for the presence of the weights file and prints a confirmation message for each model, indicating whether the file was found or not. This setup ensures that the models are properly evaluated and that the necessary weights are available for the evaluation process (Scikit-learn, 2023).

```
# Create a pandas DataFrame to organize hyperparameter configurations
# This table will display the training settings for all three models
hyperparam_df = pd.DataFrame({
    # Define the model names for easy identification
    # Used to associate each configuration with the model it belongs
    to
    "Model": ["Model A", "Model B", "Model C"],

    # Set the input image size for all models
    # 640x640 is a standard size that balances speed and accuracy
    "Image Size": [640, 640, 640],

    # Specify different optimizers to compare their effectiveness
    # AdamW, SGD, and auto represent different optimization strategies
```

```

    "Optimizer": ["AdamW", "SGD", "auto"],

    # Define the batch size for each training run
    # Smaller batches use less memory, -1 means automatic selection
    "Batch Size": ["4", "20", "-1 (Auto)"],

    # Set the initial learning rate for each model
    # Different rates test sensitivity to step size in gradient
    descent
    "Learning Rate": [0.01, 0.001, 0.0001],
})

# Display the DataFrame as a formatted table
# Shows all hyperparameter settings in an easy-to-read format
display(hyperparam_df)

# Print a footer separator line to complete the border
# Adds visual clarity at the end of the output
print("=" * 65)

# Print the table title describing the hyperparameter configurations
# Provides context for the following data table display
print("    HYPERPARAMETER CONFIGURATIONS")

# Print a header separator line for visual organization
# Separates the title from the following content
print("=" * 65)

# Creates a clear boundary for the output section
# Helps identify what information is being displayed
display(hyperparam_df)

# Print the footer separator line again
# Maintains consistency for clear, well-structured output
print("=" * 65)

```

HYPERPARAMETER CONFIGURATIONS

```

{"summary": "{\n  \"name\": \"hyperparam_df\",\n  \"rows\": 3,\n  \"fields\": [\n    {\n      \"column\": \"Model\",\n      \"properties\": {\n        \"dtype\": \"string\",\n        \"num_unique_values\": 3,\n        \"samples\": [\n          \"Model 1\",\n          \"Model 2\",\n          \"Model 3\"\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"Base Weights\",\n      \"properties\": {\n        \"dtype\": \"category\",\n        \"num_unique_values\": 1,\n        \"samples\": [\n          \"yolo26n.pt\"\n        ],\n        \"semantic_type\": \"\",

```

```
\description\: \\\"\\n    }\\n        {\\n            \\column\\\":  
\\\"Epochs\\\",\\n                \\\"properties\\\": {\\n                    \\\"dtype\\\": \\\"number\\\",\\n  
\\\"std\\\": 7,\\n                        \\\"min\\\": 25,\\n                            \\\"max\\\": 40,\\n  
\\\"num_unique_values\\\": 3,\\n                                \\\"samples\\\": [\\n                                    25\\n  
],\\n                            \\\"semantic_type\\\": \\\"\\\",\\n                                \\\"description\\\": \\\"\\\"\\n  
}\\n        },\\n        {\\n            \\\"column\\\": \\\"Image Size\\\",\\n  
\\\"properties\\\": {\\n                \\\"dtype\\\": \\\"number\\\",\\n                    \\\"std\\\":  
0,\\n                \\\"min\\\": 640,\\n                    \\\"max\\\": 640,\\n  
\\\"num_unique_values\\\": 1,\\n                    \\\"samples\\\": [\\n                        640\\n  
],\\n                \\\"semantic_type\\\": \\\"\\\",\\n                    \\\"description\\\": \\\"\\\"\\n  
}\\n        },\\n        {\\n            \\\"column\\\": \\\"Optimizer\\\",\\n  
\\\"properties\\\": {\\n                \\\"dtype\\\": \\\"string\\\",\\n  
\\\"num_unique_values\\\": 3,\\n                    \\\"samples\\\": [\\n  
\\\"AdamW\\\"\\n                        ],\\n                    \\\"semantic_type\\\": \\\"\\\",\\n  
\\\"description\\\": \\\"\\\"\\n                }\\n        },\\n        {\\n            \\\"column\\\":  
\\\"Batch Size\\\",\\n                \\\"properties\\\": {\\n                    \\\"dtype\\\":  
\\\"string\\\",\\n                        \\\"num_unique_values\\\": 3,\\n                            \\\"samples\\\":  
[\\n                                \\\"4\\\"\\n                                    ],\\n                                \\\"semantic_type\\\": \\\"\\\",\\n  
\\\"description\\\": \\\"\\\"\\n                }\\n        },\\n        {\\n            \\\"column\\\":  
\\\"Learning Rate\\\",\\n                \\\"properties\\\": {\\n                    \\\"dtype\\\":  
\\\"number\\\",\\n                        \\\"std\\\": 0.005474486277268397,\\n                            \\\"min\\\":  
0.0001,\\n                        \\\"max\\\": 0.01,\\n                            \\\"num_unique_values\\\": 3,\\n  
\\\"samples\\\": [\\n                                0.01\\n                                    ],\\n                                \\\"semantic_type\\\":  
\\\"\\\",\\n                                    \\\"description\\\": \\\"\\\"\\n                }\\n        }\\n    ]  
n}","type":"dataframe","variable name":"hyperparam df"
```

```
# Print header for Model 1 training
# This gives a visual separation before starting the training process
print("=" * 65)
```

```
# Print training configuration details for Model 1
# Displays the optimizer, epoch count, batch size, and learning rate
for clarity
print("  TRAINING MODEL 1: AdamW | 25 epochs | batch=4 | lr=0.01")
print("=" * 65)
```

```
# Initialize YOLO model with pre-trained weights
# This loads the YOLOv26 model with pre-trained weights from the
# specified file
model1 = YOLO("yolo26n.pt")
```

```
# Train model with specified hyperparameters
# Trains the model with the provided dataset, optimizer, and learning
rate
results1 = model1.train(
    data=DATA_YAML,
    epochs=25,
    imgsz=640,
    optimizer="AdamW",
    name="yolo26 model1",
```



```

    batch=4,
    lr0=0.01,
    cache="disk",
    workers=2,
    plots=True,
    verbose=True,
)

# Print completion message
# Informs the user that Model 1's training has successfully finished
print("\n🟢 Model 1 training complete!")

=====
TRAINING MODEL 1: AdamW | 25 epochs | batch=4 | lr=0.01
=====
Ultralytics 8.4.18 🐍 Python-3.12.12 torch-2.10.0+cu128 CUDA:0 (Tesla
T4, 14913MiB)
engine/trainer: agnostic_nms=False, amp=True, angle=1.0,
augment=False, auto_augment=randaugument, batch=4, bgr=0.0, box=7.5,
cache=disk, cfg=None, classes=None, close_mosaic=10, cls=0.5,
compile=False, conf=None, copy_paste=0.0, copy_paste_mode=flip,
cos_lr=False, cutmix=0.0, data=/content/dataset/data_local.yaml,
degrees=0.0, deterministic=True, device=None, dfl=1.5, dnn=False,
dropout=0.0, dynamic=False, embed=None, end2end=None, epochs=25,
erasing=0.4, exist_ok=False, fliplr=0.5, flipud=0.0,
format=torchscript, fraction=1.0, freeze=None, half=False,
hsv_h=0.015, hsv_s=0.7, hsv_v=0.4, imgsz=640, int8=False, iou=0.7,
keras=False, kobj=1.0, line_width=None, lr0=0.01, lrf=0.01,
mask_ratio=4, max_det=300, mixup=0.0, mode=train, model=yolo26n.pt,
momentum=0.937, mosaic=1.0, multi_scale=0.0, name=yolo26_model14,
nbs=64, nms=False, opset=None, optimize=False, optimizer=AdamW,
overlap_mask=True, patience=100, perspective=0.0, plots=True,
pose=12.0, pretrained=True, profile=False, project=None, rect=False,
resume=False, retina_masks=False, rle=1.0, save=True, save_conf=False,
save_crop=False, save_dir=/content/runs/detect/yolo26_model14,
save_frames=False, save_json=False, save_period=-1, save_txt=False,
scale=0.5, seed=0, shear=0.0, show=False, show_boxes=True,
show_conf=True, show_labels=True, simplify=True, single_cls=False,
source=None, split=val, stream_buffer=False, task=detect, time=None,
tracker=botsort.yaml, translate=0.1, val=True, verbose=True,
vid_stride=1, visualize=False, warmup_bias_lr=0.1, warmup_epochs=3.0,
warmup_momentum=0.8, weight_decay=0.0005, workers=2, workspace=None
Overriding model.yaml nc=80 with nc=4

      from  n    params module
arguments
   0         -1  1       464 ultralytics.nn.modules.conv.Conv
[3, 16, 3, 2]
   1         -1  1     4672 ultralytics.nn.modules.conv.Conv
[16, 32, 3, 2]

```

2	-1 1	6640	
ultralytics.nn.modules.block.C3k2			[32, 64, 1, False, 0.25]
3	-1 1	36992	ultralytics.nn.modules.conv.Conv
[64, 64, 3, 2]			
4	-1 1	26080	
ultralytics.nn.modules.block.C3k2			[64, 128, 1, False, 0.25]
5	-1 1	147712	ultralytics.nn.modules.conv.Conv
[128, 128, 3, 2]			
6	-1 1	87040	
ultralytics.nn.modules.block.C3k2			[128, 128, 1, True]
7	-1 1	295424	ultralytics.nn.modules.conv.Conv
[128, 256, 3, 2]			
8	-1 1	346112	
ultralytics.nn.modules.block.C3k2			[256, 256, 1, True]
9	-1 1	164608	
ultralytics.nn.modules.block.SPPF			[256, 256, 5, 3, True]
10	-1 1	249728	
ultralytics.nn.modules.block.C2PSA			[256, 256, 1]
11	-1 1	0	
torch.nn.modules.upsampling.Upsample			[None, 2, 'nearest']
12	[-1, 6] 1	0	
ultralytics.nn.modules.conv.Concat			[1]
13	-1 1	119808	
ultralytics.nn.modules.block.C3k2			[384, 128, 1, True]
14	-1 1	0	
torch.nn.modules.upsampling.Upsample			[None, 2, 'nearest']
15	[-1, 4] 1	0	
ultralytics.nn.modules.conv.Concat			[1]
16	-1 1	34304	
ultralytics.nn.modules.block.C3k2			[256, 64, 1, True]
17	-1 1	36992	ultralytics.nn.modules.conv.Conv
[64, 64, 3, 2]			
18	[-1, 13] 1	0	
ultralytics.nn.modules.conv.Concat			[1]
19	-1 1	95232	
ultralytics.nn.modules.block.C3k2			[192, 128, 1, True]
20	-1 1	147712	ultralytics.nn.modules.conv.Conv

```

[128, 128, 3, 2]
21          [-1, 10]  1          0
ultralytics.nn.modules.conv.Concat          [1]

22          -1  1    463104
ultralytics.nn.modules.block.C3k2          [384, 256, 1, True, 0.5,
True]
23          [16, 19, 22]  1    242736
ultralytics.nn.modules.head.Detect          [4, 1, True, [64, 128,
256]]
YOLO26n summary: 260 layers, 2,505,360 parameters, 2,505,360
gradients, 5.8 GFLOPs

Transferred 606/708 items from pretrained weights
AMP: running Automatic Mixed Precision (AMP) checks...
AMP: checks passed
train: Fast image access (ping: 0.0±0.0 ms, read: 1593.6±691.4 MB/s,
size: 61.4 KB)
train: Scanning /content/dataset/train/labels... 1676 images, 25
backgrounds, 0 corrupt: 100% _____ 1676/1676 2.4Kit/s 0.7s
train: /content/dataset/train/images/Pedestrian-Video-Data_M0V-
0015_jpg.rf.fd0d48a1835c975f77fbead6f53cb89f.jpg: 1 duplicate labels
removed
train: /content/dataset/train/images/Pedestrian-Video-Data_M0V-
0038_jpg.rf.1f960ec6897a7a1debcdcbfa8dbab64f5.jpg: 3 duplicate labels
removed
train: /content/dataset/train/images/Pedestrian-Video-Data_M0V-
0038_jpg.rf.24e0a807e680525e974cc7c9c473c3ad.jpg: 3 duplicate labels
removed
train: /content/dataset/train/images/Pedestrian-Video-Data_M0V-
0038_jpg.rf.3e8c3ca31d620a3caf30daf37e99c4a.jpg: 3 duplicate labels
removed
train: /content/dataset/train/images/Pedestrian-Video-Data_M0V-
0038_jpg.rf.a4eb66476fc41979319b5ce6171ac37b.jpg: 3 duplicate labels
removed
train: /content/dataset/train/images/Pedestrian-Video-Data_M0V-
0038_jpg.rf.f84774f9dc5eca0d927abf380b734785.jpg: 3 duplicate labels
removed
train: /content/dataset/train/images/Pedestrian-Video-Data_M0V-
0200_jpg.rf.427e230da34a884801b2499c3cdb8332.jpg: 1 duplicate labels
removed
train: /content/dataset/train/images/Pedestrian-Video-Data_M0V-
0200_jpg.rf.996fefbf4b372e68007ec0174409cb63.jpg: 1 duplicate labels
removed
train: /content/dataset/train/images/Pedestrian-Video-Data_M0V-
0200_jpg.rf.bee52b988fc94f545e36377fd89f947a.jpg: 1 duplicate labels
removed
train: /content/dataset/train/images/Pedestrian-Video-Data_M0V-
0200_jpg.rf.fd026bf67b2f6da55ce6d974f37ae485.jpg: 1 duplicate labels

```

```

removed
train: /content/dataset/train/images/Pedestrian-Video-Data_M0V-
1329_jpg.rf.8c241de3a10499b406ce863db6be890f.jpg: 1 duplicate labels
removed
train: New cache created: /content/dataset/train/labels.cache
train: Caching images (1.9GB Disk): 100% 1676/1676
172.0it/s 9.7s
augmentations: Blur(p=0.01, blur_limit=(3, 7)), MedianBlur(p=0.01,
blur_limit=(3, 7)), ToGray(p=0.01, method='weighted_average',
num_output_channels=3), CLAHE(p=0.01, clip_limit=(1.0, 4.0),
tile_grid_size=(8, 8))
val: Fast image access (ping: 0.0±0.0 ms, read: 1134.5±687.6 MB/s,
size: 64.0 KB)
val: Scanning /content/dataset/valid/labels... 465 images, 5
backgrounds, 0 corrupt: 100% 465/465 1.1Kit/s 0.4s
val: New cache created: /content/dataset/valid/labels.cache
val: Caching images (0.5GB Disk): 100% 465/465 117.4it/s
4.0s
optimizer: AdamW(lr=0.01, momentum=0.937) with parameter groups 114
weight(decay=0.0), 126 weight(decay=0.0005), 126 bias(decay=0.0)
Plotting labels to /content/runs/detect/yolo26_model14/labels.jpg...
Image sizes 640 train, 640 val
Using 2 dataloader workers
Logging results to /content/runs/detect/yolo26_model14
Starting training for 25 epochs...

```

Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
Size					
ages	Instances	Box(P	R	mAP50	mAP50-95): 100%
	59/59	4.9it/s 12.0s			
		all	465	1266	0.158 0.262
0.125	0.0474				

Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
Size					
ages	Instances	Box(P	R	mAP50	mAP50-95): 100%
	59/59	11.5it/s 5.1s			
		all	465	1266	0.22 0.346
0.168	0.0597				

Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
Size					
ages	Instances	Box(P	R	mAP50	mAP50-95): 100%
	59/59	12.2it/s 4.8s			
		all	465	1266	0.221 0.315
0.156	0.0624				

Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
Size					
ages	Instances	Box(P	R	mAP50	mAP50-95): 100%

		59/59	9.8it/s	6.0s			
		all		465	1266	0.391	0.379
0.298	0.14						

Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances		
Size	Instances	Box(P	R	mAP50	mAP50-95):	100%	
ages	59/59	10.8it/s	5.5s				
		all	465	1266	0.447	0.249	
0.259	0.114						

Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances		
Size	Instances	Box(P	R	mAP50	mAP50-95):	100%	
ages	59/59	12.2it/s	4.8s				
		all	465	1266	0.5	0.354	
0.286	0.131						

Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances		
Size	Instances	Box(P	R	mAP50	mAP50-95):	100%	
ages	59/59	11.2it/s	5.3s				
		all	465	1266	0.295	0.441	
0.321	0.155						

Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances		
Size	Instances	Box(P	R	mAP50	mAP50-95):	100%	
ages	59/59	10.0it/s	5.9s				
		all	465	1266	0.51	0.453	
0.346	0.159						

Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances		
Size	Instances	Box(P	R	mAP50	mAP50-95):	100%	
ages	59/59	12.4it/s	4.7s				
		all	465	1266	0.296	0.45	
0.315	0.131						

Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances		
Size	Instances	Box(P	R	mAP50	mAP50-95):	100%	
ages	59/59	12.1it/s	4.9s				
		all	465	1266	0.408	0.502	
0.43	0.217						

Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances		
Size	Instances	Box(P	R	mAP50	mAP50-95):	100%	
ages	59/59	10.1it/s	5.8s				

0.448	0.222	all	465	1266	0.396	0.549
-------	-------	-----	-----	------	-------	-------

Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
Size					
ages	Instances	Box(P	R	mAP50	mAP50-95): 100%
59/59	12.4it/s	4.8s			

0.448	0.236	all	465	1266	0.372	0.535
-------	-------	-----	-----	------	-------	-------

Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
Size					
ages	Instances	Box(P	R	mAP50	mAP50-95): 100%
59/59	10.9it/s	5.4s			

0.473	0.256	all	465	1266	0.444	0.562
-------	-------	-----	-----	------	-------	-------

Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
Size					
ages	Instances	Box(P	R	mAP50	mAP50-95): 100%
59/59	9.7it/s	6.1s			

0.481	0.232	all	465	1266	0.432	0.546
-------	-------	-----	-----	------	-------	-------

Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
Size					
ages	Instances	Box(P	R	mAP50	mAP50-95): 100%
59/59	12.4it/s	4.7s			

0.457	0.222	all	465	1266	0.451	0.524
-------	-------	-----	-----	------	-------	-------

Closing dataloader mosaic
albumentations: Blur(p=0.01, blur_limit=(3, 7)), MedianBlur(p=0.01, blur_limit=(3, 7)), ToGray(p=0.01, method='weighted_average', num_output_channels=3), CLAHE(p=0.01, clip_limit=(1.0, 4.0), tile_grid_size=(8, 8))

Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
Size					
ages	Instances	Box(P	R	mAP50	mAP50-95): 100%
59/59	12.0it/s	4.9s			

0.473	0.237	all	465	1266	0.421	0.556
-------	-------	-----	-----	------	-------	-------

Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
Size					
ages	Instances	Box(P	R	mAP50	mAP50-95): 100%
59/59	11.2it/s	5.3s			

0.491	0.265	all	465	1266	0.41	0.566
-------	-------	-----	-----	------	------	-------

Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
ages	Instances	Box(P	R	mAP50	mAP50-95):	100%
	59/59	9.4it/s	6.3s			
0.495	0.27	all	465	1266	0.45	0.582

Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
ages	Instances	Box(P	R	mAP50	mAP50-95):	100%
	59/59	11.0it/s	5.4s			
0.529	0.297	all	465	1266	0.48	0.588

Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
ages	Instances	Box(P	R	mAP50	mAP50-95):	100%
	59/59	12.0it/s	4.9s			
0.504	0.267	all	465	1266	0.479	0.576

Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
ages	Instances	Box(P	R	mAP50	mAP50-95):	100%
	59/59	12.2it/s	4.8s			
0.517	0.295	all	465	1266	0.423	0.618

Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
ages	Instances	Box(P	R	mAP50	mAP50-95):	100%
	59/59	9.5it/s	6.2s			
0.534	0.308	all	465	1266	0.453	0.61

Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
ages	Instances	Box(P	R	mAP50	mAP50-95):	100%
	59/59	10.0it/s	5.9s			
0.532	0.303	all	465	1266	0.512	0.587

Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
ages	Instances	Box(P	R	mAP50	mAP50-95):	100%
	59/59	12.2it/s	4.8s			
0.538	0.314	all	465	1266	0.495	0.604

	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
--	-------	---------	----------	----------	-----------	-----------

Size	ages	Instances	Box(P	R	mAP50	mAP50-95):	100%
		59/59	12.2it/s	4.8s			
			all	465	1266	0.512	0.593
0.542		0.318					

25 epochs completed in 0.531 hours.

Optimizer stripped from

/content/runs/detect/yolo26_model14/weights/last.pt, 5.4MB

Optimizer stripped from

/content/runs/detect/yolo26_model14/weights/best.pt, 5.4MB

Validating /content/runs/detect/yolo26_model14/weights/best.pt...

Ultralytics 8.4.18 Python-3.12.12 torch-2.10.0+cu128 CUDA:0 (Tesla T4, 14913MiB)

YOLO26n summary (fused): 122 layers, 2,375,616 parameters, 0 gradients, 5.2 GFLOPs

Size	ages	Instances	Box(P	R	mAP50	mAP50-95):	100%
		59/59	14.3it/s	4.1s			
			all	465	1266	0.513	0.593
0.543		0.318					
			crossing	139	272	0.496	0.654
0.515		0.298					
			jaywalking	243	331	0.668	0.831
0.832		0.517					
			pedestrian	133	396	0.381	0.194
0.253		0.137					
			waiting	152	267	0.506	0.693
0.572		0.322					

Speed: 0.4ms preprocess, 3.8ms inference, 0.0ms loss, 0.3ms postprocess per image

Results saved to /content/runs/detect/yolo26_model14

Model 1 training complete!

Delete model and free memory

Removes model and results variables to free up memory

`del model1, results1`

Trigger garbage collection to reclaim unused memory

Ensures that Python's memory management system frees up memory properly

`gc.collect()`

Clear the GPU cache to release unused memory

This helps prevent memory issues when using CUDA for deep learning tasks

`torch.cuda.empty_cache()`

Print memory clearance message


```

# Informs the user that memory has been successfully cleared
print("Memory cleared")

Memory cleared

# Print header for Model 2 training
# This provides a clear separation before starting the training for Model 2
print("=" * 65)

# Print training configuration details for Model 2
# Displays the optimizer, epoch count, batch size, and learning rate for clarity
print("  TRAINING MODEL 2: SGD | 30 epochs | batch=20 | lr=0.001")
print("=" * 65)

# Initialize YOLO model with pre-trained weights
# Loads the YOLOv26 model with pre-trained weights for use in training
model2 = YOLO("yolo26n.pt")

# Train Model 2 with different hyperparameters and save results
# This initiates training with a different optimizer and settings for Model 2
results2 = model2.train(
    data=DATA_YAML,
    epochs=30,
    imgsz=640,
    optimizer="SGD",
    name="yolo26_model2",
    batch=20,
    lr0=0.001,
    cache="disk",
    workers=2,
    plots=True,
    verbose=True,
)

# Print completion message
# Informs the user that Model 2's training has been successfully completed
print("\n Model 2 training complete!")

=====
  TRAINING MODEL 2: SGD | 30 epochs | batch=20 | lr=0.001
=====
Ultralytics 8.4.18 Python-3.12.12 torch-2.10.0+cu128 CUDA:0 (Tesla T4, 14913MiB)
engine/trainer: agnostic_nms=False, amp=True, angle=1.0,
augment=False, auto_augment=randaugment, batch=20, bgr=0.0, box=7.5,
cache=disk, cfg=None, classes=None, close_mosaic=10, cls=0.5,

```

```

compile=False, conf=None, copy_paste=0.0, copy_paste_mode=flip,
cos_lr=False, cutmix=0.0, data=/content/dataset/data_local.yaml,
degrees=0.0, deterministic=True, device=None, dfl=1.5, dnn=False,
dropout=0.0, dynamic=False, embed=None, end2end=None, epochs=30,
erasing=0.4, exist_ok=False, flipplr=0.5, flipud=0.0,
format=torchscript, fraction=1.0, freeze=None, half=False,
hsv_h=0.015, hsv_s=0.7, hsv_v=0.4, imgsz=640, int8=False, iou=0.7,
keras=False, kobj=1.0, line_width=None, lr0=0.001, lrf=0.01,
mask_ratio=4, max_det=300, mixup=0.0, mode=train, model=yolo26n.pt,
momentum=0.937, mosaic=1.0, multi_scale=0.0, name=yolo26_model25,
nbs=64, nms=False, opset=None, optimize=False, optimizer=SGD,
overlap_mask=True, patience=100, perspective=0.0, plots=True,
pose=12.0, pretrained=True, profile=False, project=None, rect=False,
resume=False, retina_masks=False, rle=1.0, save=True, save_conf=False,
save_crop=False, save_dir=/content/runs/detect/yolo26_model25,
save_frames=False, save_json=False, save_period=-1, save_txt=False,
scale=0.5, seed=0, shear=0.0, show=False, show_boxes=True,
show_conf=True, show_labels=True, simplify=True, single_cls=False,
source=None, split=val, stream_buffer=False, task=detect, time=None,
tracker=botsort.yaml, translate=0.1, val=True, verbose=True,
vid_stride=1, visualize=False, warmup_bias_lr=0.1, warmup_epochs=3.0,
warmup_momentum=0.8, weight_decay=0.0005, workers=2, workspace=None
Overriding model.yaml nc=80 with nc=4

```

	from	n	params	module
arguments				
0	-1	1	464	ultralytics.nn.modules.conv.Conv
[3, 16, 3, 2]				
1	-1	1	4672	ultralytics.nn.modules.conv.Conv
[16, 32, 3, 2]				
2	-1	1	6640	
ultralytics.nn.modules.block.C3k2				[32, 64, 1, False, 0.25]
3	-1	1	36992	ultralytics.nn.modules.conv.Conv
[64, 64, 3, 2]				
4	-1	1	26080	
ultralytics.nn.modules.block.C3k2				[64, 128, 1, False, 0.25]
5	-1	1	147712	ultralytics.nn.modules.conv.Conv
[128, 128, 3, 2]				
6	-1	1	87040	
ultralytics.nn.modules.block.C3k2				[128, 128, 1, True]
7	-1	1	295424	ultralytics.nn.modules.conv.Conv
[128, 256, 3, 2]				
8	-1	1	346112	
ultralytics.nn.modules.block.C3k2				[256, 256, 1, True]
9	-1	1	164608	
ultralytics.nn.modules.block.SPPF				[256, 256, 5, 3, True]

```

10          -1  1    249728
ultralytics.nn.modules.block.C2PSA          [256, 256, 1]

11          -1  1         0
torch.nn.modules.upsampling.Upsample        [None, 2, 'nearest']

12         [-1, 6]  1         0
ultralytics.nn.modules.conv.Concat          [1]

13          -1  1    119808
ultralytics.nn.modules.block.C3k2          [384, 128, 1, True]

14          -1  1         0
torch.nn.modules.upsampling.Upsample        [None, 2, 'nearest']

15         [-1, 4]  1         0
ultralytics.nn.modules.conv.Concat          [1]

16          -1  1    34304
ultralytics.nn.modules.block.C3k2          [256, 64, 1, True]

17          -1  1    36992  ultralytics.nn.modules.conv.Conv
[64, 64, 3, 2]

18         [-1, 13]  1         0
ultralytics.nn.modules.conv.Concat          [1]

19          -1  1    95232
ultralytics.nn.modules.block.C3k2          [192, 128, 1, True]

20          -1  1    147712  ultralytics.nn.modules.conv.Conv
[128, 128, 3, 2]

21         [-1, 10]  1         0
ultralytics.nn.modules.conv.Concat          [1]

22          -1  1    463104
ultralytics.nn.modules.block.C3k2          [384, 256, 1, True, 0.5,
True]

23         [16, 19, 22]  1    242736
ultralytics.nn.modules.head.Detect          [4, 1, True, [64, 128,
256]]
YOL026n summary: 260 layers, 2,505,360 parameters, 2,505,360
gradients, 5.8 GFLOPs

```

Transferred 606/708 items from pretrained weights

AMP: running Automatic Mixed Precision (AMP) checks...

AMP: checks passed ☐

train: Fast image access ☐ (ping: 0.0±0.0 ms, read: 1309.2±755.2 MB/s, size: 63.8 KB)

train: Scanning /content/dataset/train/labels.cache... 1676 images, 25

```
backgrounds, 0 corrupt: 100% ————— 1676/1676 639.1Mit/s 0.0s
train: /content/dataset/train/images/Pedestrian-Video-Data_M0V-
0015_jpg.rf.fd0d48a1835c975f77fbead6f53cb89f.jpg: 1 duplicate labels
removed
train: /content/dataset/train/images/Pedestrian-Video-Data_M0V-
0038_jpg.rf.1f960ec6897a7a1debcd bfa8dbab64f5.jpg: 3 duplicate labels
removed
train: /content/dataset/train/images/Pedestrian-Video-Data_M0V-
0038_jpg.rf.24e0a807e680525e974cc7c9c473c3ad.jpg: 3 duplicate labels
removed
train: /content/dataset/train/images/Pedestrian-Video-Data_M0V-
0038_jpg.rf.3e8c3ca31d620a3caf30dafa37e99c4a.jpg: 3 duplicate labels
removed
train: /content/dataset/train/images/Pedestrian-Video-Data_M0V-
0038_jpg.rf.a4eb66476fc41979319b5ce6171ac37b.jpg: 3 duplicate labels
removed
train: /content/dataset/train/images/Pedestrian-Video-Data_M0V-
0038_jpg.rf.f84774f9dc5eca0d927abf380b734785.jpg: 3 duplicate labels
removed
train: /content/dataset/train/images/Pedestrian-Video-Data_M0V-
0200_jpg.rf.427e230da34a884801b2499c3cdb8332.jpg: 1 duplicate labels
removed
train: /content/dataset/train/images/Pedestrian-Video-Data_M0V-
0200_jpg.rf.996fefbf4b372e68007ec0174409cb63.jpg: 1 duplicate labels
removed
train: /content/dataset/train/images/Pedestrian-Video-Data_M0V-
0200_jpg.rf.bee52b988fc94f545e36377fd89f947a.jpg: 1 duplicate labels
removed
train: /content/dataset/train/images/Pedestrian-Video-Data_M0V-
0200_jpg.rf.fd026bf67b2f6da55ce6d974f37ae485.jpg: 1 duplicate labels
removed
train: /content/dataset/train/images/Pedestrian-Video-Data_M0V-
1329_jpg.rf.8c241de3a10499b406ce863db6be890f.jpg: 1 duplicate labels
removed
train: Caching images (1.9GB Disk): 100% ————— 1676/1676
41.4Kit/s 0.0s
augmentations: Blur(p=0.01, blur_limit=(3, 7)), MedianBlur(p=0.01,
blur_limit=(3, 7)), ToGray(p=0.01, method='weighted_average',
num_output_channels=3), CLAHE(p=0.01, clip_limit=(1.0, 4.0),
tile_grid_size=(8, 8))
val: Fast image access □ (ping: 0.0±0.0 ms, read: 597.9±469.5 MB/s,
size: 63.5 KB)
val: Scanning /content/dataset/valid/labels.cache... 465 images, 5
backgrounds, 0 corrupt: 100% ————— 465/465 19.5Mit/s 0.0s
val: Caching images (0.5GB Disk): 100% ————— 465/465 15.4Kit/s
0.0s
optimizer: SGD(lr=0.001, momentum=0.937) with parameter groups 114
weight(decay=0.0), 126 weight(decay=0.00046875), 126 bias(decay=0.0)
Plotting labels to /content/runs/detect/yolo26_model25/labels.jpg...
```

Image sizes 640 train, 640 val
Using 2 dataloader workers
Logging results to /content/runs/detect/yolo26_model25
Starting training for 30 epochs...

Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
Size					
ages	Instances	Box(P	R	mAP50	mAP50-95): 100%
	12/12	1.3it/s 9.3s			
		all	465	1266	0.000346 0.0346
0.000221	7.91e-05				

Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
Size					
ages	Instances	Box(P	R	mAP50	mAP50-95): 100%
	12/12	3.8it/s 3.1s			
		all	465	1266	0.784 0.0234
0.0761	0.0376				

Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
Size					
ages	Instances	Box(P	R	mAP50	mAP50-95): 100%
	12/12	3.9it/s 3.1s			
		all	465	1266	0.143 0.219
0.113	0.0596				

Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
Size					
ages	Instances	Box(P	R	mAP50	mAP50-95): 100%
	12/12	3.1it/s 3.8s			
		all	465	1266	0.211 0.333
0.173	0.0888				

Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
Size					
ages	Instances	Box(P	R	mAP50	mAP50-95): 100%
	12/12	3.5it/s 3.5s			
		all	465	1266	0.25 0.315
0.216	0.112				

Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
Size					
ages	Instances	Box(P	R	mAP50	mAP50-95): 100%
	12/12	3.8it/s 3.2s			
		all	465	1266	0.271 0.364
0.253	0.137				

Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
Size					
ages	Instances	Box(P	R	mAP50	mAP50-95): 100%

0.291	0.155	12/12	3.2it/s	3.8s	all	465	1266	0.316	0.395
-------	-------	-------	---------	------	-----	-----	------	-------	-------

Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
Size	ages	Instances	Box(P	R	mAP50 mAP50-95): 100%
0.296	0.164	12/12	3.9it/s	3.1s	all
			465	1266	0.309 0.443

Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
Size	ages	Instances	Box(P	R	mAP50 mAP50-95): 100%
0.324	0.171	12/12	3.8it/s	3.1s	all
			465	1266	0.351 0.408

Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
Size	ages	Instances	Box(P	R	mAP50 mAP50-95): 100%
0.337	0.179	12/12	3.1it/s	3.8s	all
			465	1266	0.362 0.427

Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
Size	ages	Instances	Box(P	R	mAP50 mAP50-95): 100%
0.333	0.189	12/12	3.6it/s	3.4s	all
			465	1266	0.315 0.44

Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
Size	ages	Instances	Box(P	R	mAP50 mAP50-95): 100%
0.328	0.181	12/12	3.8it/s	3.1s	all
			465	1266	0.32 0.452

Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
Size	ages	Instances	Box(P	R	mAP50 mAP50-95): 100%
0.357	0.209	12/12	3.7it/s	3.2s	all
			465	1266	0.352 0.462

Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
Size	ages	Instances	Box(P	R	mAP50 mAP50-95): 100%
		12/12	3.0it/s	3.9s	

0.35	0.201	all	465	1266	0.333	0.473
------	-------	-----	-----	------	-------	-------

Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
Size					
ages	Instances	Box(P	R	mAP50	mAP50-95): 100%
	12/12	3.9it/s	3.1s		
0.377	0.216	all	465	1266	0.382 0.457

Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
Size					
ages	Instances	Box(P	R	mAP50	mAP50-95): 100%
	12/12	3.9it/s	3.1s		
0.376	0.213	all	465	1266	0.372 0.509

Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
Size					
ages	Instances	Box(P	R	mAP50	mAP50-95): 100%
	12/12	2.9it/s	4.2s		
0.387	0.225	all	465	1266	0.379 0.5

Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
Size					
ages	Instances	Box(P	R	mAP50	mAP50-95): 100%
	12/12	3.9it/s	3.1s		
0.381	0.222	all	465	1266	0.347 0.519

Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
Size					
ages	Instances	Box(P	R	mAP50	mAP50-95): 100%
	12/12	3.9it/s	3.1s		
0.398	0.236	all	465	1266	0.386 0.493

Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
Size					
ages	Instances	Box(P	R	mAP50	mAP50-95): 100%
	12/12	3.5it/s	3.5s		
0.394	0.227	all	465	1266	0.373 0.516

Closing dataloader mosaic
albumentations: Blur(p=0.01, blur_limit=(3, 7)), MedianBlur(p=0.01, blur_limit=(3, 7)), ToGray(p=0.01, method='weighted_average', num_output_channels=3), CLAHE(p=0.01, clip_limit=(1.0, 4.0), tile_grid_size=(8, 8))

Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
ages	Instances	Box(P	R	mAP50	mAP50-95):	100%
	12/12	3.8it/s	3.2s			
0.393	0.222	all	465	1266	0.398	0.481

Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
ages	Instances	Box(P	R	mAP50	mAP50-95):	100%
	12/12	3.8it/s	3.1s			
0.412	0.237	all	465	1266	0.381	0.525

Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
ages	Instances	Box(P	R	mAP50	mAP50-95):	100%
	12/12	3.9it/s	3.0s			
0.416	0.243	all	465	1266	0.391	0.542

Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
ages	Instances	Box(P	R	mAP50	mAP50-95):	100%
	12/12	3.3it/s	3.6s			
0.416	0.242	all	465	1266	0.408	0.51

Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
ages	Instances	Box(P	R	mAP50	mAP50-95):	100%
	12/12	2.8it/s	4.3s			
0.424	0.245	all	465	1266	0.409	0.511

Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
ages	Instances	Box(P	R	mAP50	mAP50-95):	100%
	12/12	3.6it/s	3.3s			
0.434	0.25	all	465	1266	0.434	0.525

Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
ages	Instances	Box(P	R	mAP50	mAP50-95):	100%
	12/12	3.9it/s	3.1s			
0.431	0.25	all	465	1266	0.411	0.523

	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
--	-------	---------	----------	----------	-----------	-----------

Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
ages	12/12	3.9it/s	3.1s			
0.434	0.254	all	465	1266	0.408	0.54

Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
ages	12/12	3.9it/s	3.1s			
0.43	0.25	all	465	1266	0.408	0.523

Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
ages	12/12	3.2it/s	3.8s			
0.429	0.246	all	465	1266	0.422	0.516

30 epochs completed in 0.289 hours.

Optimizer stripped from

/content/runs/detect/yolo26_model25/weights/last.pt, 5.4MB

Optimizer stripped from

/content/runs/detect/yolo26_model25/weights/best.pt, 5.4MB

Validating /content/runs/detect/yolo26_model25/weights/best.pt...

Ultralytics 8.4.18 Python-3.12.12 torch-2.10.0+cu128 CUDA:0 (Tesla T4, 14913MiB)

YOLO26n summary (fused): 122 layers, 2,375,616 parameters, 0 gradients, 5.2 GFLOPs

ages	Instances	Box(P	R	mAP50	mAP50-95):
0.434	0.255	all	465	1266	0.408 0.541
0.398	0.237	crossing	139	272	0.376 0.614
0.733	0.467	jaywalking	243	331	0.606 0.752
0.206	0.107	pedestrian	133	396	0.308 0.197
0.4	0.207	waiting	152	267	0.34 0.603

Speed: 0.2ms preprocess, 2.4ms inference, 0.0ms loss, 0.8ms postprocess per image

Results saved to /content/runs/detect/yolo26_model25

Model 2 training complete!

```

# Delete model and free memory
# Removes model and results variables to free up memory
del model2, results2

# Trigger garbage collection to reclaim unused memory
# Ensures that Python's memory management system frees up memory properly
gc.collect()

# Clear the GPU cache to release unused memory
# Prevents potential memory issues when working with CUDA for deep learning tasks
torch.cuda.empty_cache()

# Print memory clearance message
# Notifies the user that memory has been successfully cleared
print("Memory cleared")

Memory cleared

# Print header for Model 3 training
# This creates a visual separation before starting the training for Model 3
print("=" * 65)

# Print training configuration details for Model 3
# Displays the optimizer, epoch count, batch size, and learning rate for clarity
print("  TRAINING MODEL 3: auto | 40 epochs | batch=-1 | lr=0.0001")
print("=" * 65)

# Initialize YOLO model with pre-trained weights
# Loads the YOLOv26 model with pre-trained weights for use in training
model3 = YOLO("yolo26n.pt")

# Train model with automatic optimizer and batch size
# This starts the training process with automatic batch sizing and optimizer
results3 = model3.train(
    data=DATA_YAML,
    epochs=40,
    imgsz=640,
    optimizer="auto",
    name="yolo26_model3",
    batch=-1,
    lr0=0.0001,
    cache="disk",
    workers=2,
    plots=True,
    verbose=True,

```

```

)

# Print completion message
# Informs the user that Model 3's training has been successfully
completed
print("\n Model 3 training complete!")

=====
TRAINING MODEL 3: auto | 40 epochs | batch=-1 | lr=0.0001
=====
Ultralytics 8.4.18 Python-3.12.12 torch-2.10.0+cu128 CUDA:0 (Tesla
T4, 14913MiB)
engine/trainer: agnostic_nms=False, amp=True, angle=1.0,
augment=False, auto_augment=randaugument, batch=-1, bgr=0.0, box=7.5,
cache=disk, cfg=None, classes=None, close_mosaic=10, cls=0.5,
compile=False, conf=None, copy_paste=0.0, copy_paste_mode=flip,
cos_lr=False, cutmix=0.0, data=/content/dataset/data_local.yaml,
degrees=0.0, deterministic=True, device=None, dfl=1.5, dnn=False,
dropout=0.0, dynamic=False, embed=None, end2end=None, epochs=40,
erasing=0.4, exist_ok=False, flipplr=0.5, flipud=0.0,
format=torchscript, fraction=1.0, freeze=None, half=False,
hsv_h=0.015, hsv_s=0.7, hsv_v=0.4, imgsz=640, int8=False, iou=0.7,
keras=False, kobj=1.0, line_width=None, lr0=0.0001, lrf=0.01,
mask_ratio=4, max_det=300, mixup=0.0, mode=train, model=yolo26n.pt,
momentum=0.937, mosaic=1.0, multi_scale=0.0, name=yolo26_model3,
nbs=64, nms=False, opset=None, optimize=False, optimizer=auto,
overlap_mask=True, patience=100, perspective=0.0, plots=True,
pose=12.0, pretrained=True, profile=False, project=None, rect=False,
resume=False, retina_masks=False, rle=1.0, save=True, save_conf=False,
save_crop=False, save_dir=/content/runs/detect/yolo26_model3,
save_frames=False, save_json=False, save_period=-1, save_txt=False,
scale=0.5, seed=0, shear=0.0, show=False, show_boxes=True,
show_conf=True, show_labels=True, simplify=True, single_cls=False,
source=None, split=val, stream_buffer=False, task=detect, time=None,
tracker=botsort.yaml, translate=0.1, val=True, verbose=True,
vid_stride=1, visualize=False, warmup_bias_lr=0.1, warmup_epochs=3.0,
warmup_momentum=0.8, weight_decay=0.0005, workers=2, workspace=None
Overriding model.yaml nc=80 with nc=4

      from  n  params module
arguments
  0          -1  1      464 ultralytics.nn.modules.conv.Conv
[3, 16, 3, 2]
  1          -1  1     4672 ultralytics.nn.modules.conv.Conv
[16, 32, 3, 2]
  2          -1  1     6640
ultralytics.nn.modules.block.C3k2          [32, 64, 1, False, 0.25]
  3          -1  1    36992 ultralytics.nn.modules.conv.Conv
[64, 64, 3, 2]

```

4	-1 1	26080	
ultralytics.nn.modules.block.C3k2			[64, 128, 1, False, 0.25]
5	-1 1	147712	ultralytics.nn.modules.conv.Conv
[128, 128, 3, 2]			
6	-1 1	87040	
ultralytics.nn.modules.block.C3k2			[128, 128, 1, True]
7	-1 1	295424	ultralytics.nn.modules.conv.Conv
[128, 256, 3, 2]			
8	-1 1	346112	
ultralytics.nn.modules.block.C3k2			[256, 256, 1, True]
9	-1 1	164608	
ultralytics.nn.modules.block.SPPF			[256, 256, 5, 3, True]
10	-1 1	249728	
ultralytics.nn.modules.block.C2PSA			[256, 256, 1]
11	-1 1	0	
torch.nn.modules.upsampling.Upsample			[None, 2, 'nearest']
12	[-1, 6] 1	0	
ultralytics.nn.modules.conv.Concat			[1]
13	-1 1	119808	
ultralytics.nn.modules.block.C3k2			[384, 128, 1, True]
14	-1 1	0	
torch.nn.modules.upsampling.Upsample			[None, 2, 'nearest']
15	[-1, 4] 1	0	
ultralytics.nn.modules.conv.Concat			[1]
16	-1 1	34304	
ultralytics.nn.modules.block.C3k2			[256, 64, 1, True]
17	-1 1	36992	ultralytics.nn.modules.conv.Conv
[64, 64, 3, 2]			
18	[-1, 13] 1	0	
ultralytics.nn.modules.conv.Concat			[1]
19	-1 1	95232	
ultralytics.nn.modules.block.C3k2			[192, 128, 1, True]
20	-1 1	147712	ultralytics.nn.modules.conv.Conv
[128, 128, 3, 2]			
21	[-1, 10] 1	0	
ultralytics.nn.modules.conv.Concat			[1]
22	-1 1	463104	

ultralytics.nn.modules.block.C3k2 [384, 256, 1, True, 0.5, True]

23 [16, 19, 22] 1 242736

ultralytics.nn.modules.head.Detect [4, 1, True, [64, 128, 256]]

YOLO26n summary: 260 layers, 2,505,360 parameters, 2,505,360 gradients, 5.8 GFLOPs

Transferred 606/708 items from pretrained weights

AMP: running Automatic Mixed Precision (AMP) checks...

AMP: checks passed

train: Fast image access (ping: 0.0±0.0 ms, read: 1315.6±505.7 MB/s, size: 63.8 KB)

train: Scanning /content/dataset/train/labels.cache... 1676 images, 25 backgrounds, 0 corrupt: 100% 1676/1676 439.4Mit/s 0.0s

train: /content/dataset/train/images/Pedestrian-Video-Data_MOV-0015_jpg.rf.fd0d48a1835c975f77fbeat6f53cb89f.jpg: 1 duplicate labels removed

train: /content/dataset/train/images/Pedestrian-Video-Data_MOV-0038_jpg.rf.1f960ec6897a7a1debcd bfa8dbab64f5.jpg: 3 duplicate labels removed

train: /content/dataset/train/images/Pedestrian-Video-Data_MOV-0038_jpg.rf.24e0a807e680525e974cc7c9c473c3ad.jpg: 3 duplicate labels removed

train: /content/dataset/train/images/Pedestrian-Video-Data_MOV-0038_jpg.rf.3e8c3ca31d620a3caf30dafa37e99c4a.jpg: 3 duplicate labels removed

train: /content/dataset/train/images/Pedestrian-Video-Data_MOV-0038_jpg.rf.a4eb66476fc41979319b5ce6171ac37b.jpg: 3 duplicate labels removed

train: /content/dataset/train/images/Pedestrian-Video-Data_MOV-0038_jpg.rf.f84774f9dc5eca0d927abf380b734785.jpg: 3 duplicate labels removed

train: /content/dataset/train/images/Pedestrian-Video-Data_MOV-0200_jpg.rf.427e230da34a884801b2499c3cdb8332.jpg: 1 duplicate labels removed

train: /content/dataset/train/images/Pedestrian-Video-Data_MOV-0200_jpg.rf.996fefbf4b372e68007ec0174409cb63.jpg: 1 duplicate labels removed

train: /content/dataset/train/images/Pedestrian-Video-Data_MOV-0200_jpg.rf.bee52b988fc94f545e36377fd89f947a.jpg: 1 duplicate labels removed

train: /content/dataset/train/images/Pedestrian-Video-Data_MOV-0200_jpg.rf.fd026bf67b2f6da55ce6d974f37ae485.jpg: 1 duplicate labels removed

train: /content/dataset/train/images/Pedestrian-Video-Data_MOV-1329_jpg.rf.8c241de3a10499b406ce863db6be890f.jpg: 1 duplicate labels removed

augmentations: Blur(p=0.01, blur_limit=(3, 7)), MedianBlur(p=0.01,

```

blur_limit=(3, 7)), ToGray(p=0.01, method='weighted_average',
num_output_channels=3), CLAHE(p=0.01, clip_limit=(1.0, 4.0),
tile_grid_size=(8, 8))
AutoBatch: Computing optimal batch size for imgsz=640 at 60.0% CUDA
memory utilization.
AutoBatch: CUDA:0 (Tesla T4) 14.56G total, 0.13G reserved, 0.09G
allocated, 14.35G free

```

	Params	GFLOPs	GPU_mem (GB)	forward (ms)	backward (ms)
input		output			
	2505360	5.778	0.545	44.55	nan
(1, 3, 640, 640)			list		
	2505360	11.56	1.053	45.61	nan
(2, 3, 640, 640)			list		
	2505360	23.11	1.804	54.38	nan
(4, 3, 640, 640)			list		
	2505360	46.23	3.355	58.22	nan
(8, 3, 640, 640)			list		
	2505360	92.45	6.321	74.65	nan
(16, 3, 640, 640)			list		

```

AutoBatch: Using batch-size 22 for CUDA:0 8.86G/14.56G (61%)
train: Fast image access (ping: 0.0±0.0 ms, read: 1361.7±408.6 MB/s,
size: 58.5 KB)
train: Scanning /content/dataset/train/labels.cache... 1676 images, 25
backgrounds, 0 corrupt: 100% 1676/1676 390.5Mit/s 0.0s
train: /content/dataset/train/images/Pedestrian-Video-Data_M0V-
0015_jpg.rf.fd0d48a1835c975f77fbead6f53cb89f.jpg: 1 duplicate labels
removed
train: /content/dataset/train/images/Pedestrian-Video-Data_M0V-
0038_jpg.rf.1f960ec6897a7a1debcd bfa8dbab64f5.jpg: 3 duplicate labels
removed
train: /content/dataset/train/images/Pedestrian-Video-Data_M0V-
0038_jpg.rf.24e0a807e680525e974cc7c9c473c3ad.jpg: 3 duplicate labels
removed
train: /content/dataset/train/images/Pedestrian-Video-Data_M0V-
0038_jpg.rf.3e8c3ca31d620a3caf30dafa37e99c4a.jpg: 3 duplicate labels
removed
train: /content/dataset/train/images/Pedestrian-Video-Data_M0V-
0038_jpg.rf.a4eb66476fc41979319b5ce6171ac37b.jpg: 3 duplicate labels
removed
train: /content/dataset/train/images/Pedestrian-Video-Data_M0V-
0038_jpg.rf.f84774f9dc5eca0d927abf380b734785.jpg: 3 duplicate labels
removed
train: /content/dataset/train/images/Pedestrian-Video-Data_M0V-
0200_jpg.rf.427e230da34a884801b2499c3cdb8332.jpg: 1 duplicate labels
removed
train: /content/dataset/train/images/Pedestrian-Video-Data_M0V-
0200_jpg.rf.996fefbf4b372e68007ec0174409cb63.jpg: 1 duplicate labels
removed
train: /content/dataset/train/images/Pedestrian-Video-Data_M0V-

```

```

0200_jpg.rf.bee52b988fc94f545e36377fd89f947a.jpg: 1 duplicate labels removed
train: /content/dataset/train/images/Pedestrian-Video-Data_M0V-0200_jpg.rf.fd026bf67b2f6da55ce6d974f37ae485.jpg: 1 duplicate labels removed
train: /content/dataset/train/images/Pedestrian-Video-Data_M0V-1329_jpg.rf.8c241de3a10499b406ce863db6be890f.jpg: 1 duplicate labels removed
train: Caching images (1.9GB Disk): 100% _____ 1676/1676 36.7Kit/s 0.0s
augmentations: Blur(p=0.01, blur_limit=(3, 7)), MedianBlur(p=0.01, blur_limit=(3, 7)), ToGray(p=0.01, method='weighted_average', num_output_channels=3), CLAHE(p=0.01, clip_limit=(1.0, 4.0), tile_grid_size=(8, 8))
val: Fast image access □ (ping: 0.0±0.0 ms, read: 21.1±19.4 MB/s, size: 64.7 KB)
val: Scanning /content/dataset/valid/labels.cache... 465 images, 5 backgrounds, 0 corrupt: 100% _____ 465/465 17.3Mit/s 0.0s
val: Caching images (0.5GB Disk): 100% _____ 465/465 15.0Kit/s 0.0s
optimizer: 'optimizer=auto' found, ignoring 'lr0=0.0001' and 'momentum=0.937' and determining best 'optimizer', 'lr0' and 'momentum' automatically...
optimizer: AdamW(lr=0.00125, momentum=0.9) with parameter groups 114 weight(decay=0.0), 126 weight(decay=0.000515625), 126 bias(decay=0.0)
Plotting labels to /content/runs/detect/yolo26_model3/labels.jpg...
Image sizes 640 train, 640 val
Using 2 dataloader workers
Logging results to /content/runs/detect/yolo26_model3
Starting training for 40 epochs...

```

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances
Size					
ages	Instances	Box(P	R	mAP50	mAP50-95): 100%
_____	11/11	1.9it/s	5.8s		
	all	465	1266	0.00586	0.442
0.113	0.0546				

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances
Size					
ages	Instances	Box(P	R	mAP50	mAP50-95): 100%
_____	11/11	3.0it/s	3.7s		
	all	465	1266	0.36	0.356
0.31	0.155				

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances
Size					
ages	Instances	Box(P	R	mAP50	mAP50-95): 100%
_____	11/11	2.9it/s	3.8s		
	all	465	1266	0.383	0.452

0.379	0.2						
Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	
ages	Instances	Box(P	R	mAP50	mAP50-95):	100%	
	11/11	3.5it/s	3.1s				
		all	465	1266	0.337	0.488	
0.392	0.194						
Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	
ages	Instances	Box(P	R	mAP50	mAP50-95):	100%	
	11/11	3.6it/s	3.0s				
		all	465	1266	0.576	0.428	
0.397	0.203						
Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	
ages	Instances	Box(P	R	mAP50	mAP50-95):	100%	
	11/11	3.0it/s	3.6s				
		all	465	1266	0.417	0.478	
0.408	0.219						
Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	
ages	Instances	Box(P	R	mAP50	mAP50-95):	100%	
	11/11	3.0it/s	3.6s				
		all	465	1266	0.415	0.502	
0.439	0.22						
Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	
ages	Instances	Box(P	R	mAP50	mAP50-95):	100%	
	11/11	3.6it/s	3.1s				
		all	465	1266	0.448	0.555	
0.461	0.242						
Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	
ages	Instances	Box(P	R	mAP50	mAP50-95):	100%	
	11/11	3.6it/s	3.1s				
		all	465	1266	0.484	0.578	
0.516	0.283						
Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	
ages	Instances	Box(P	R	mAP50	mAP50-95):	100%	
	11/11	2.8it/s	3.9s				
		all	465	1266	0.498	0.575	
0.497	0.289						

Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
ages	Instances		Box(P	R	mAP50	mAP50-95): 100%
	11/11	3.1it/s	3.5s			
		all	465	1266	0.503	0.596
0.547	0.325					

Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
ages	Instances		Box(P	R	mAP50	mAP50-95): 100%
	11/11	3.6it/s	3.1s			
		all	465	1266	0.526	0.615
0.555	0.336					

Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
ages	Instances		Box(P	R	mAP50	mAP50-95): 100%
	11/11	3.6it/s	3.0s			
		all	465	1266	0.505	0.588
0.562	0.333					

Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
ages	Instances		Box(P	R	mAP50	mAP50-95): 100%
	11/11	2.8it/s	4.0s			
		all	465	1266	0.528	0.648
0.554	0.319					

Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
ages	Instances		Box(P	R	mAP50	mAP50-95): 100%
	11/11	3.0it/s	3.6s			
		all	465	1266	0.523	0.649
0.572	0.328					

Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
ages	Instances		Box(P	R	mAP50	mAP50-95): 100%
	11/11	3.5it/s	3.1s			
		all	465	1266	0.547	0.575
0.577	0.339					

Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
ages	Instances		Box(P	R	mAP50	mAP50-95): 100%
	11/11	3.7it/s	3.0s			
		all	465	1266	0.557	0.627
0.617	0.364					

Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
ages	Instances	Box(P	R	mAP50	mAP50-95):	100%
	11/11	2.7it/s	4.1s			
		all	465	1266	0.571	0.635
0.625	0.377					
Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
ages	Instances	Box(P	R	mAP50	mAP50-95):	100%
	11/11	3.6it/s	3.1s			
		all	465	1266	0.532	0.592
0.594	0.371					
Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
ages	Instances	Box(P	R	mAP50	mAP50-95):	100%
	11/11	3.4it/s	3.3s			
		all	465	1266	0.602	0.645
0.643	0.395					
Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
ages	Instances	Box(P	R	mAP50	mAP50-95):	100%
	11/11	2.5it/s	4.3s			
		all	465	1266	0.587	0.622
0.62	0.377					
Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
ages	Instances	Box(P	R	mAP50	mAP50-95):	100%
	11/11	3.6it/s	3.1s			
		all	465	1266	0.573	0.671
0.613	0.368					
Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
ages	Instances	Box(P	R	mAP50	mAP50-95):	100%
	11/11	3.5it/s	3.1s			
		all	465	1266	0.586	0.652
0.642	0.379					
Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
ages	Instances	Box(P	R	mAP50	mAP50-95):	100%
	11/11	3.0it/s	3.7s			
		all	465	1266	0.585	0.682
0.633	0.402					
	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances

Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
ages	Instances	Box(P	R	mAP50	mAP50-95):	100%
11/11	3.4it/s	3.3s				
all	465	1266	0.635	0.653		
0.659	0.416					

Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
ages	Instances	Box(P	R	mAP50	mAP50-95):	100%
11/11	3.5it/s	3.1s				
all	465	1266	0.611	0.661		
0.656	0.412					

Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
ages	Instances	Box(P	R	mAP50	mAP50-95):	100%
11/11	3.4it/s	3.2s				
all	465	1266	0.602	0.669		
0.658	0.41					

Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
ages	Instances	Box(P	R	mAP50	mAP50-95):	100%
11/11	2.8it/s	4.0s				
all	465	1266	0.6	0.669		
0.66	0.408					

Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
ages	Instances	Box(P	R	mAP50	mAP50-95):	100%
11/11	3.7it/s	3.0s				
all	465	1266	0.599	0.668		
0.673	0.42					

Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
ages	Instances	Box(P	R	mAP50	mAP50-95):	100%
11/11	3.6it/s	3.0s				
all	465	1266	0.619	0.654		
0.667	0.428					

Closing dataloader mosaic
 albumentations: Blur(p=0.01, blur_limit=(3, 7)), MedianBlur(p=0.01, blur_limit=(3, 7)), ToGray(p=0.01, method='weighted_average', num_output_channels=3), CLAHE(p=0.01, clip_limit=(1.0, 4.0), tile_grid_size=(8, 8))

Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
ages	Instances	Box(P	R	mAP50	mAP50-95):	100%
11/11	2.5it/s	4.4s				

0.673	0.419	all	465	1266	0.633	0.642
Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
ages	Instances	Box(P	R	mAP50	mAP50-95):	100%
	11/11	3.3it/s	3.4s			
0.666	0.42	all	465	1266	0.626	0.676
Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
ages	Instances	Box(P	R	mAP50	mAP50-95):	100%
	11/11	3.6it/s	3.0s			
0.69	0.432	all	465	1266	0.643	0.677
Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
ages	Instances	Box(P	R	mAP50	mAP50-95):	100%
	11/11	3.6it/s	3.1s			
0.687	0.435	all	465	1266	0.65	0.674
Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
ages	Instances	Box(P	R	mAP50	mAP50-95):	100%
	11/11	3.7it/s	3.0s			
0.691	0.444	all	465	1266	0.665	0.67
Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
ages	Instances	Box(P	R	mAP50	mAP50-95):	100%
	11/11	3.0it/s	3.7s			
0.689	0.433	all	465	1266	0.656	0.671
Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
ages	Instances	Box(P	R	mAP50	mAP50-95):	100%
	11/11	2.8it/s	4.0s			
0.698	0.443	all	465	1266	0.661	0.683
Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
ages	Instances	Box(P	R	mAP50	mAP50-95):	100%
	11/11	3.5it/s	3.2s			
		all	465	1266	0.676	0.658

```

0.699      0.447

      Epoch      GPU_mem    box_loss    cls_loss    dfl_loss    Instances
Size
ages  Instances      Box(P          R      mAP50  mAP50-95): 100%
----- 11/11 3.6it/s 3.1s
      all          465          1266          0.674          0.679
0.697      0.445

      Epoch      GPU_mem    box_loss    cls_loss    dfl_loss    Instances
Size
ages  Instances      Box(P          R      mAP50  mAP50-95): 100%
----- 11/11 3.5it/s 3.1s
      all          465          1266          0.673          0.675
0.698      0.445

40 epochs completed in 0.378 hours.
Optimizer stripped from
/content/runs/detect/yolo26_model3/weights/last.pt, 5.4MB
Optimizer stripped from
/content/runs/detect/yolo26_model3/weights/best.pt, 5.4MB

Validating /content/runs/detect/yolo26_model3/weights/best.pt...
Ultralytics 8.4.18 Python-3.12.12 torch-2.10.0+cu128 CUDA:0 (Tesla
T4, 14913MiB)
YOLO26n summary (fused): 122 layers, 2,375,616 parameters, 0
gradients, 5.2 GFLOPs
ages  Instances      Box(P          R      mAP50  mAP50-95): 100%
----- 11/11 1.8it/s 6.1s
      all          465          1266          0.676          0.656
0.699      0.447
      crossing          139          272          0.679          0.717
0.736      0.473
      jaywalking          243          331          0.851          0.863
0.905      0.641
      pedestrian          133          396          0.515          0.361
0.454      0.256
      waiting          152          267          0.658          0.684
0.702      0.42
Speed: 0.2ms preprocess, 2.6ms inference, 0.0ms loss, 1.0ms
postprocess per image
Results saved to /content/runs/detect/yolo26_model3

Model 3 training complete!

```

Section 3 - Model Evaluation (The Confusion Matrix)

In this cell, the notebook iterates over each model's training directory to display its corresponding training loss and metrics curves. It first checks if the results.png image, which

contains the training curves, exists in the specified directory. If the image is found, it uses Matplotlib to load and display the training curve with an appropriately sized figure, ensuring no clipping by adjusting the layout. The image is displayed with the title reflecting the model's name for clarity. If the image is not found, a warning message is printed, indicating the missing results file. This section provides a visual representation of the model's performance over training, helping assess its progress and metrics (Matplotlib Documentation, 2023).

```
# Define dataset configuration path
# Specifies the location of the YAML configuration file for the dataset
DATA_YAML = "/content/dataset/data_local.yaml"

# Map model names to their output directories
# Associates each model with the respective directory where training results are stored
train_dirs = {
    "Model 1": "runs/detect/yolo26_model14",
    "Model 2": "runs/detect/yolo26_model25",
    "Model 3": "runs/detect/yolo26_model3",
}

# Verify training outputs exist
# Iterates through each model's output directory to check for the existence of the best model weights
for name, dir_path in train_dirs.items():
    best = os.path.join(dir_path, "weights", "best.pt")

    # Check if the best model weights file exists
    # Informs the user whether the model weights are found or not
    if os.path.exists(best):
        print(f"□ {name}: {dir_path}")
    else:
        print(f"□ {name}: NOT FOUND at {dir_path}")

□ Model 1: runs/detect/yolo26_model14
□ Model 2: runs/detect/yolo26_model25
□ Model 3: runs/detect/yolo26_model3
```

Hyperparameter Table

```
# Create a comprehensive DataFrame with all hyperparameter settings
# This provides a complete reference table for model configurations
hyperparam_df = pd.DataFrame({
    # Model names for identification
    # Labels each row with the corresponding model
    "Model": ["Model 1", "Model 2", "Model 3"],

    # Number of training epochs for each model
    # Shows the variation in training duration
    "Epochs": [25, 30, 40],
```

```

# Input image dimensions for all models
# All models use the same image size for consistency
"Image Size": [640, 640, 640],

# Optimizer algorithm used for each model
display(hyperparam_df)
# Compares AdamW, SGD, and automatic selection# Shows all settings
in an organized, readable format

"Optimizer": ["AdamW", "SGD", "auto"],# Display the hyperparameter
table

# Batch size configuration for each model
print("=" * 70)
# Shows range from small fixed to automatic selection# Completes
the visual frame around the title

"Batch Size": ["4", "20", "-1 (auto)"],# Print footer separator
line

# Initial learning rate for each model
print("  HYPERPARAMETER
SETTINGS – ALL 3 MODELS")
# Demonstrates different learning rate strategies# Describes the
comprehensive hyperparameter information

"Learning Rate": [0.01, 0.001, 0.0001],# Print the table title

# Number of data loading workers
print("=" * 70)
# All models use 2 workers for parallel data loading# Creates
visual boundary for the table

"Workers": [2, 2, 2],

# Caching strategy for dataset}
# All models use disk caching for faster data access      "Cache":
["disk", "disk", "disk"]

=====
HYPERPARAMETER SETTINGS – ALL 3 MODELS
=====

{"summary":{"\n  \"name\": \"hyperparam_df\", \n  \"rows\": 3, \n
\"fields\": [\n    {\n      \"column\": \"Model\", \n
\"properties\": {\n        \"dtype\": \"string\", \n
\"num_unique_values\": 3, \n        \"samples\": [\n
\"Model 1\", \n        \"Model 2\", \n        \"Model 3\" \n
], \n
\"semantic_type\": \"\", \n        \"description\": \"\" \n
} \n
n    }, \n    {\n      \"column\": \"Epochs\", \n
\"properties\":

```

```
{\n      \"dtype\": \"number\", \n      \"std\": 7, \n      \"min\": 25, \n      \"max\": 40, \n      \"num_unique_values\": 3, \n      \"samples\": [\n        25, \n        30, \n        40\n      ], \n      \"semantic_type\": \"\", \n      \"description\": \"\"\n    }, \n    {\n      \"column\": \"Image Size\", \n      \"properties\": {\n        \"dtype\": \"number\", \n        \"std\": 0, \n        \"min\": 640, \n        \"max\": 640, \n        \"num_unique_values\": 1, \n        \"samples\": [\n          640\n        ], \n        \"semantic_type\": \"\", \n        \"description\": \"\"\n      }, \n      \"column\": \"Optimizer\", \n      \"properties\": {\n        \"dtype\": \"string\", \n        \"num_unique_values\": 3, \n        \"samples\": [\n          \"AdamW\", \n          \"\", \n          \"\n        ], \n        \"semantic_type\": \"\", \n        \"description\": \"\"\n      }, \n      \"column\": \"Batch Size\", \n      \"properties\": {\n        \"dtype\": \"string\", \n        \"num_unique_values\": 3, \n        \"samples\": [\n          \"4\", \n          \"\", \n          \"\n        ], \n        \"semantic_type\": \"\", \n        \"description\": \"\"\n      }, \n      \"column\": \"Learning Rate\", \n      \"properties\": {\n        \"dtype\": \"number\", \n        \"std\": 0.005474486277268397, \n        \"min\": 0.0001, \n        \"max\": 0.01, \n        \"num_unique_values\": 3, \n        \"samples\": [\n          0.01, \n          \"\", \n          \"\n        ], \n        \"semantic_type\": \"\", \n        \"description\": \"\"\n      }, \n      \"column\": \"Workers\", \n      \"properties\": {\n        \"dtype\": \"number\", \n        \"std\": 0, \n        \"min\": 2, \n        \"max\": 2, \n        \"num_unique_values\": 1, \n        \"samples\": [\n          2\n        ], \n        \"semantic_type\": \"\", \n        \"description\": \"\"\n      }, \n      \"column\": \"Cache\", \n      \"properties\": {\n        \"dtype\": \"category\", \n        \"num_unique_values\": 1, \n        \"samples\": [\n          \"disk\", \n          \"\n        ], \n        \"semantic_type\": \"\", \n        \"description\": \"\"\n      }\n    }\n  ], \n  \"type\": \"dataframe\", \n  \"variable_name\": \"hyperparam_df\"}
```

Training Curves

```
# Display training curves for each model
# Iterates over each model to display its training loss and metrics
# curves
for name, dir_path in train_dirs.items():
    results_img = os.path.join(dir_path, \"results.png\")

    # Check if the results image exists in the specified directory
    # If the image exists, it will be displayed; otherwise, a warning
    # is printed
    if os.path.exists(results_img):
        print(f\"\\n{'='*65}\")
        print(f\"    {name} - Training Loss & Metrics Curves\")
        print(f\"{'='*65}\")
```

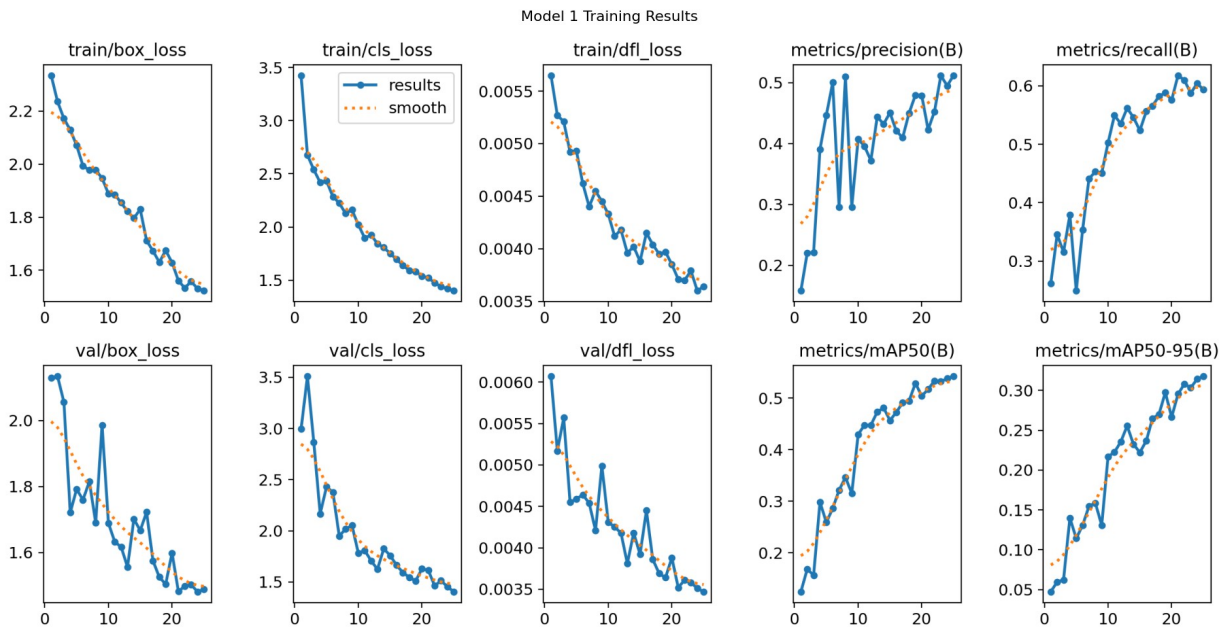


```

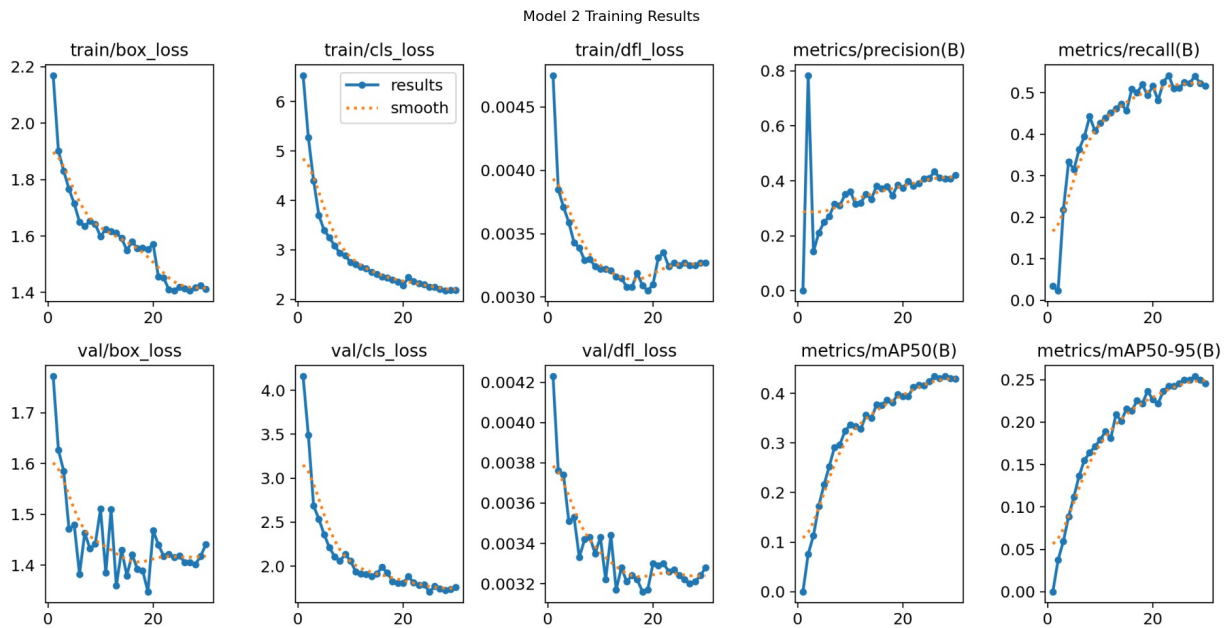
# Load and display the results image
# Uses Matplotlib to load and display the training curve image
img = mpimg.imread(results_img)
fig, ax = plt.subplots(1, 1, figsize=(20, 10))
ax.imshow(img)
ax.set_title(f"{name} Training Results", fontsize=16)
ax.axis("off") # Turn off the axis for a clean presentation
plt.tight_layout() # Adjust layout to ensure no clipping
plt.show() # Display the image
else:
    print(f"⚠ {name}: results.png not found")

```

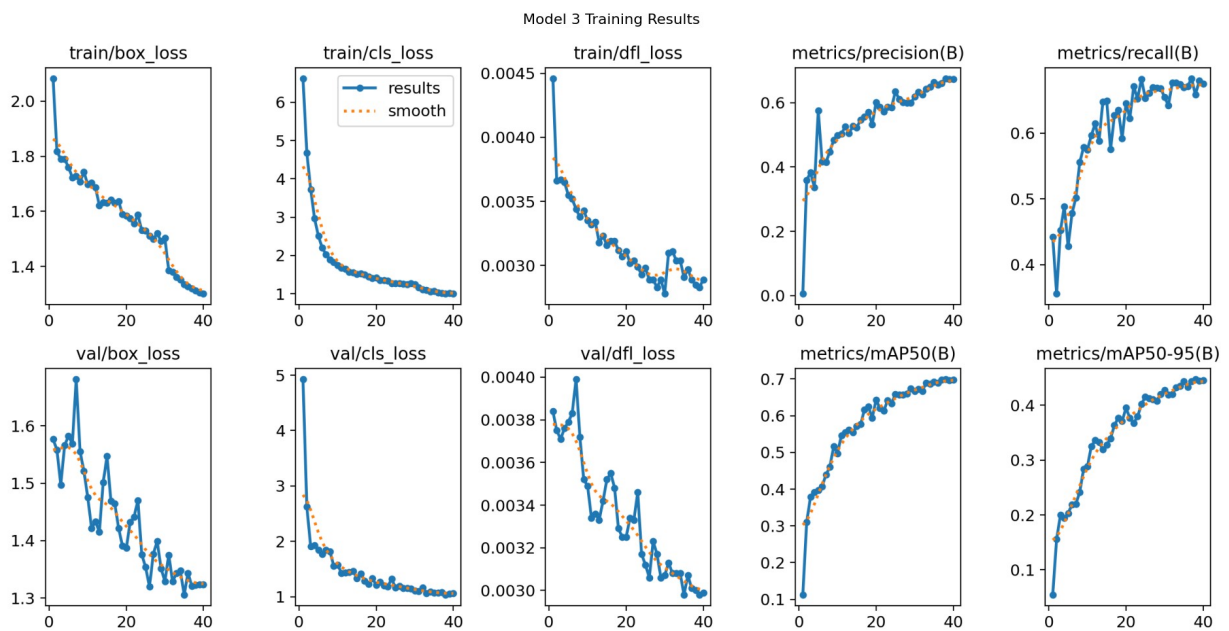
Model 1 – Training Loss & Metrics Curves



Model 2 – Training Loss & Metrics Curves



Model 3 – Training Loss & Metrics Curves



Confusion Matrices Side-by-Side

This cell is dedicated to comparing the confusion matrices of all trained models side-by-side for easy visual comparison. It begins by initializing a plot with one row and three columns to display the confusion matrices for each model. For each model, the code checks if the normalized

confusion matrix image exists; if not, it uses the raw confusion matrix image instead. If the file exists, it is loaded and displayed in the corresponding subplot, with the model name set as the title of the subplot. If the image is missing, a "Not Found" message is displayed in the subplot. The section concludes by setting a title for the entire figure and adjusting the layout to ensure that the matrices are displayed clearly without any overlapping. This comparison helps evaluate the models' performance and identify potential areas of improvement (Matplotlib Documentation, 2023).

```
# Create side-by-side comparison of confusion matrices
# Initializes a plot with 1 row and 3 columns to display confusion
matrices for all models
fig, axes = plt.subplots(1, 3, figsize=(24, 8))

# Iterate through each model to load and display its confusion matrix
# For each model, checks if the normalized confusion matrix exists,
otherwise uses the raw one
for idx, (name, dir_path) in enumerate(train_dirs.items()):
    cm_norm = os.path.join(dir_path,
"confusion_matrix_normalized.png")
    cm_raw = os.path.join(dir_path, "confusion_matrix.png")
    chosen = cm_norm if os.path.exists(cm_norm) else cm_raw

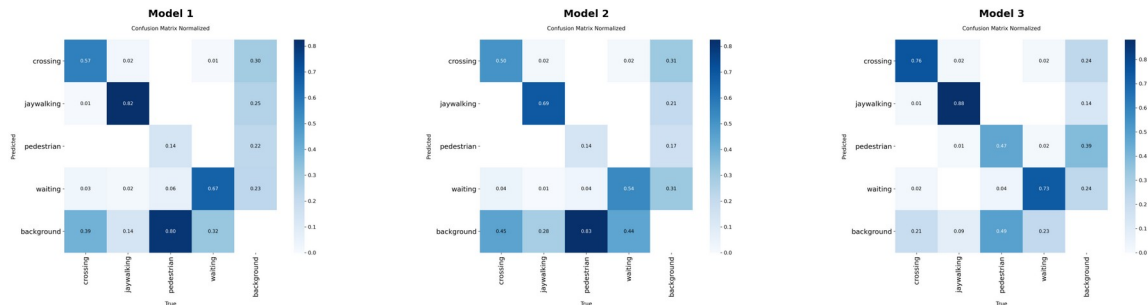
    # Check if the chosen confusion matrix file exists
    # If it exists, display it; otherwise, show a "Not Found" message
    on the plot
    if os.path.exists(chosen):
        img = mpimg.imread(chosen)
        axes[idx].imshow(img)
        axes[idx].set_title(f"{name}", fontsize=14, fontweight="bold")
        axes[idx].axis("off") # Turn off the axis for clean
presentation
    else:
        # Display "Not Found" text in the subplot if the confusion
matrix is missing
        axes[idx].text(
            0.5, 0.5, f"{name}\nNot Found",
            ha="center", va="center", fontsize=14,
        )
        axes[idx].axis("off") # Turn off the axis for a cleaner view

# Set the title for the entire figure
# Adds a title for the comparison of confusion matrices across all
models
plt.suptitle(
    "Confusion Matrices – All 3 Models",
    fontsize=18, fontweight="bold", y=1.02,
)

# Adjust the layout to ensure no overlap and display the plot
```

```
plt.tight_layout()
plt.show() # Show the figure with confusion matrices
```

Confusion Matrices — All 3 Models

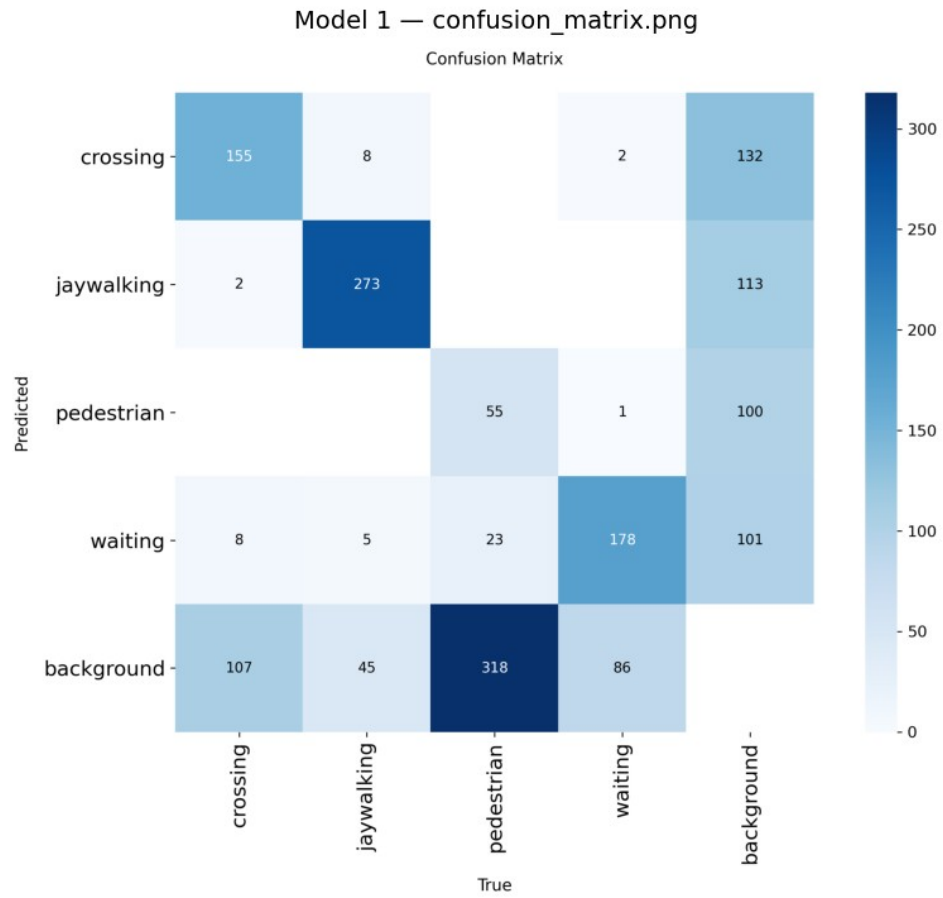


```
# Display individual confusion matrices for each model
# Iterates over each model and its confusion matrix files for display
for name, dir_path in train_dirs.items():
    for cm_file in [
        "confusion_matrix.png",
        "confusion_matrix_normalized.png",
    ]:
        cm_path = os.path.join(dir_path, cm_file)

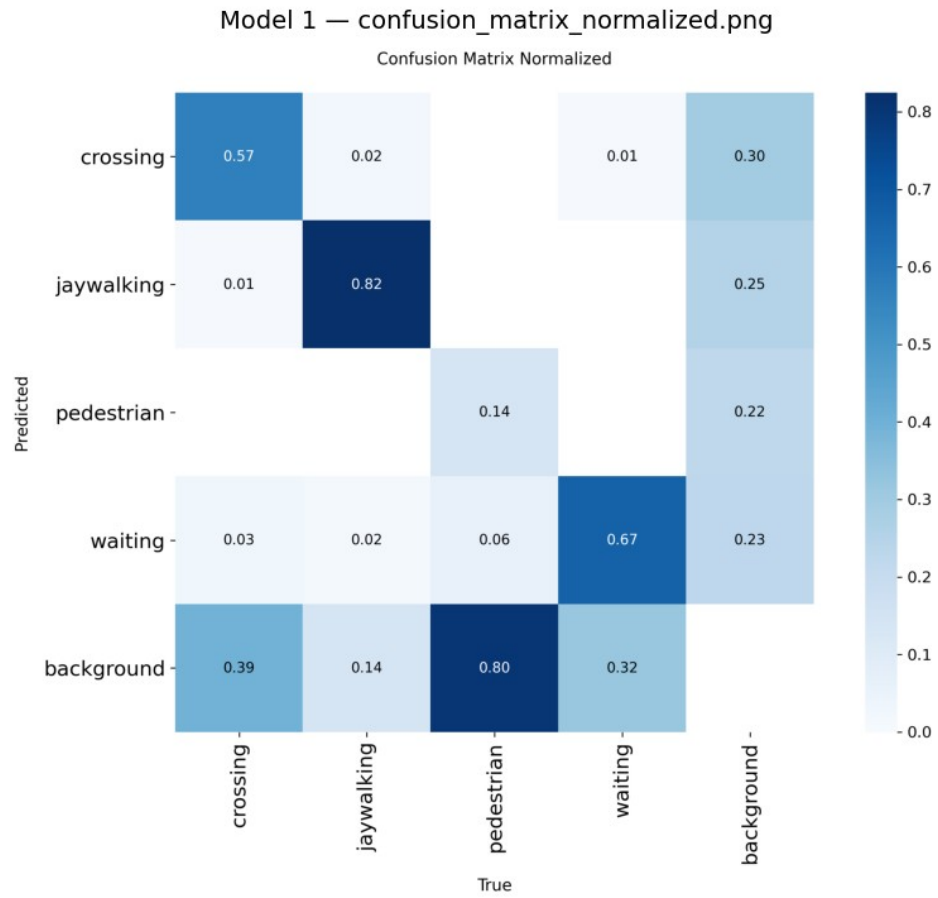
        # Check if the confusion matrix file exists
        # If the file exists, it will be loaded and displayed
        if os.path.exists(cm_path):
            print(f"\n{name} - {cm_file}")

            # Load and display the confusion matrix
            # Uses Matplotlib to read and display the confusion matrix
            image
            img = mpimg.imread(cm_path)
            fig, ax = plt.subplots(1, 1, figsize=(10, 10))
            ax.imshow(img)
            ax.set_title(f"{name} - {cm_file}", fontsize=14) # Sets
the title for the image
            ax.axis("off") # Hides the axis for a cleaner
presentation
            plt.tight_layout() # Ensures the layout is neat and no
clipping occurs
            plt.show() # Displays the confusion matrix image
```

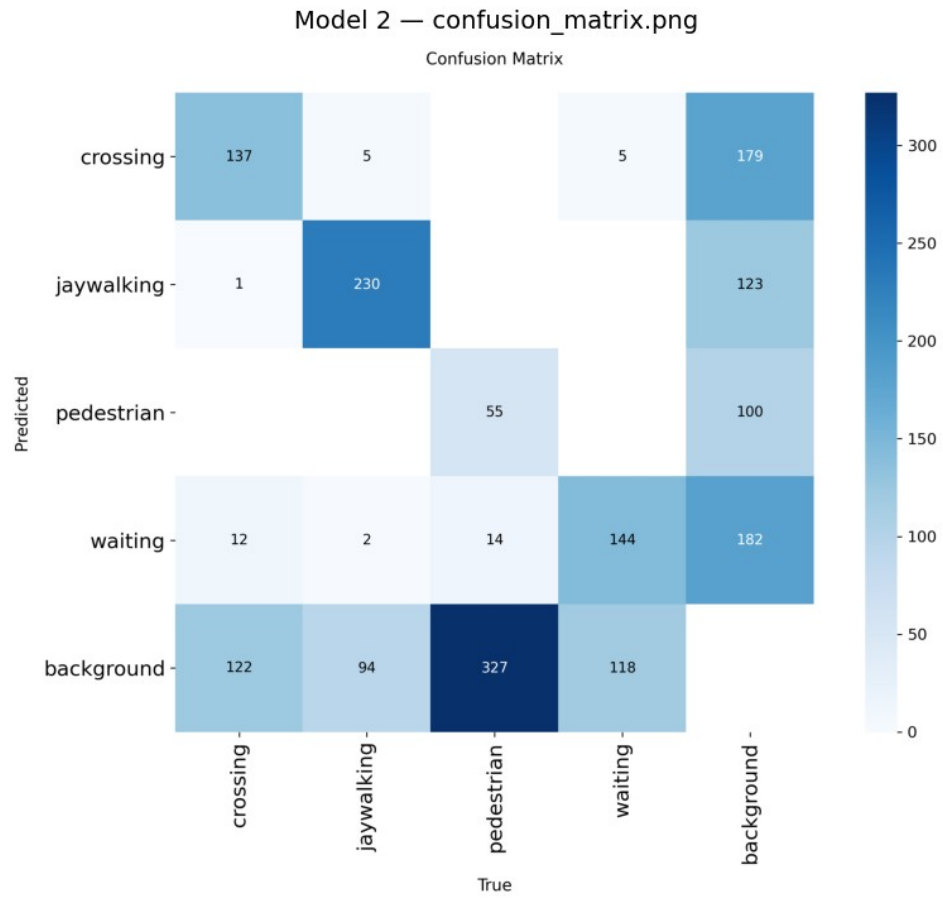
Model 1 – confusion_matrix.png



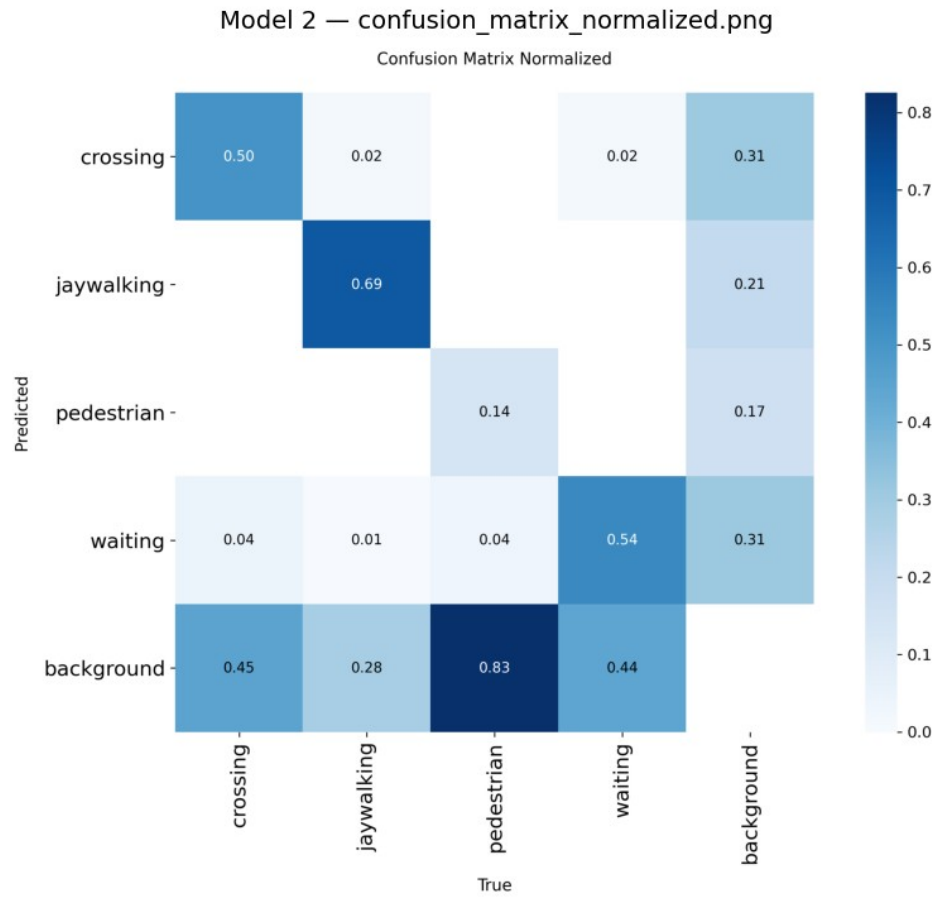
Model 1 — confusion_matrix_normalized.png



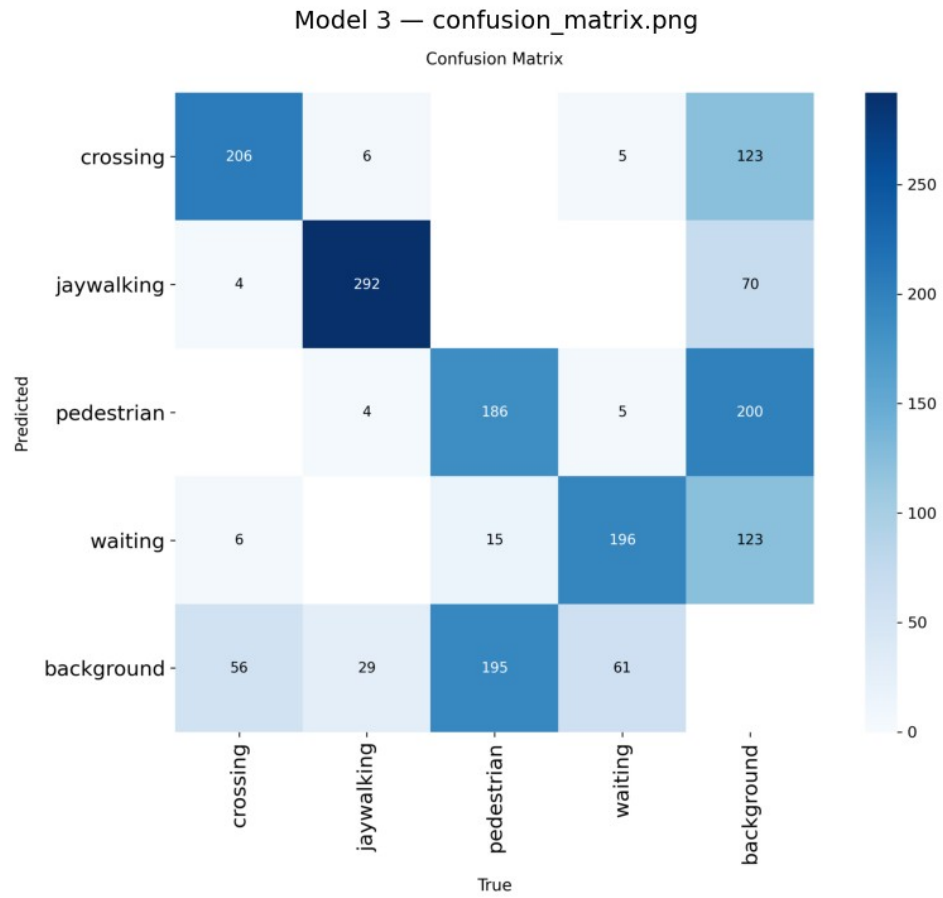
Model 2 — confusion_matrix.png



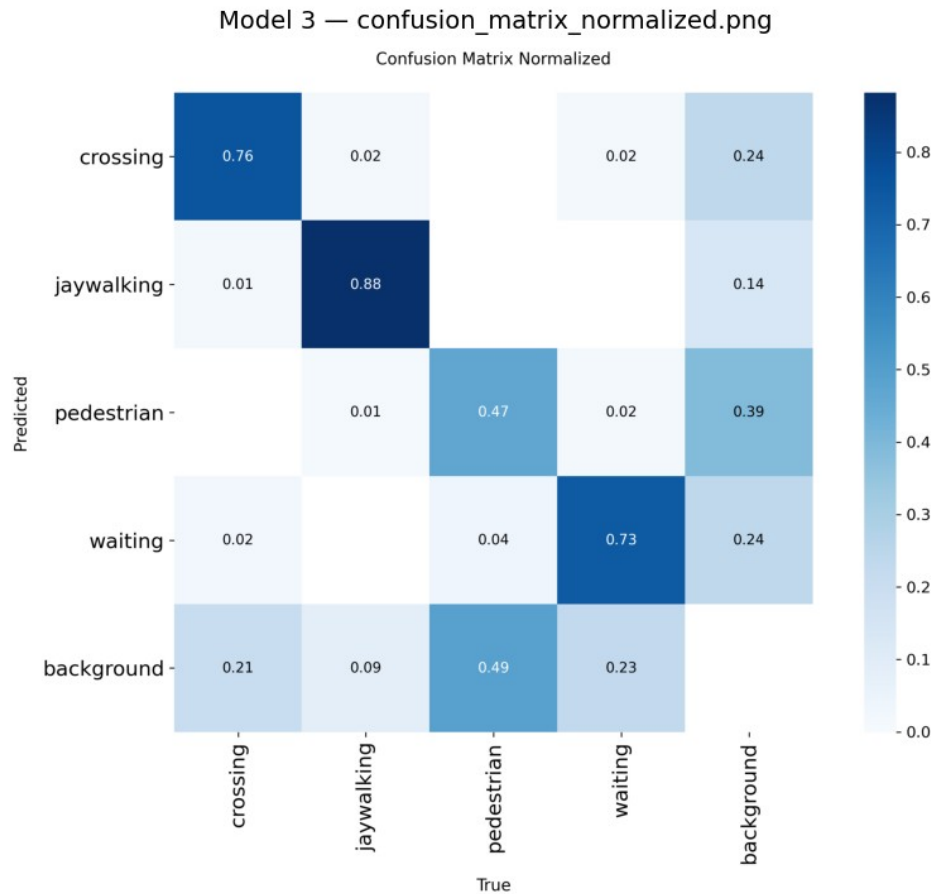
Model 2 — confusion_matrix_normalized.png



Model 3 — confusion_matrix.png



Model 3 — confusion_matrix_normalized.png



Section 4 - Generating the Performance Table

This section focuses on validating each trained model and generating a performance table to store and display the results. It starts by importing the necessary libraries, such as torch for model handling and gc for memory management. The dataset's class names are defined, corresponding to the object types the model is trained to detect. The section then iterates over each model's training directory to check if the best.pt weights file exists. If the file is found, the model validation process proceeds, and the results are stored in a dictionary for later use. The performance metrics for each model are organized, ensuring that validation results are properly handled and stored for easy comparison in the following sections. This setup provides a structured way to collect and analyze model performance (PyTorch Documentation, 2023).

```
# Run validation for each trained model
# Import necessary libraries: torch for model handling and gc for
memory management
import torch, gc

# Define class names for the dataset
# These are the object classes the model is trained to detect
class_names = ["crossing", "jaywalking", "pedestrian", "waiting"]
```

```

val_results = {} # Dictionary to store validation results for each
model

# Iterate through each model and perform validation
# For each model, check if the best weights file exists before running
validation
for name, dir_path in train_dirs.items():
    best_weights = os.path.join(dir_path, "weights", "best.pt")

    # Check if the best weights file exists
    # If the file exists, proceed with model validation
    if os.path.exists(best_weights):
        print(f"\n{'='*65}")
        print(f"    Validating {name}")
        print(f"{'='*65}")

        # Load the model with the best weights
        # Initialize the model with pre-trained weights for evaluation
        model = YOLO(best_weights)
        metrics = model.val(data=DATA_YAML, imgsz=640, verbose=True)

        # Extract overall metrics for validation
        # Collect precision, recall, mAP50, and F1 score for the
model's performance
        precision = metrics.box.mp
        recall = metrics.box.mr
        map50 = metrics.box.map50
        map50_95 = metrics.box.map
        f1 = (
            2 * (precision * recall) / (precision + recall)
            if (precision + recall) > 0
            else 0.0
        )

        # Store overall metrics in the results dictionary
        # Round and store mAP50, mAP50-95, Precision, Recall, and F1
Score for each model
        val_results[name] = {
            "overall": {
                "mAP50": round(map50, 4),
                "mAP50-95": round(map50_95, 4),
                "Precision": round(precision, 4),
                "Recall": round(recall, 4),
                "F1 Score": round(f1, 4),
            },
            "per_class": [], # List to store per-class metrics
        }

        # Extract per-class metrics
        # Loop through the class names and extract Precision, Recall,

```

```

and AP50 for each class
    p_arr = metrics.box.p
    r_arr = metrics.box.r
    ap50_arr = metrics.box.ap50

    for i, cls_name in enumerate(class_names):
        p = p_arr[i] if i < len(p_arr) else 0
        r = r_arr[i] if i < len(r_arr) else 0
        ap50 = ap50_arr[i] if i < len(ap50_arr) else 0
        cls_f1 = 2 * (p * r) / (p + r) if (p + r) > 0 else 0

        # Store per-class metrics in the results dictionary
        val_results[name]["per_class"].append({
            "Class": cls_name,
            "AP50": round(ap50, 4),
            "Precision": round(p, 4),
            "Recall": round(r, 4),
            "F1 Score": round(cls_f1, 4),
        })

    # Clean up memory after each model validation
    # Delete model and metrics objects to free up memory and
prevent overload
    del model, metrics
    gc.collect() # Force garbage collection
    torch.cuda.empty_cache() # Clear GPU cache to free memory
    print(f" {name} cleared from memory")
else:
    # If the best weights file is not found, print a warning
    print(f" {name}: best.pt not found at {best_weights}")

```

Validating Model 1

```

=====
Ultralytics 8.4.18 □ Python-3.12.12 torch-2.10.0+cu128 CUDA:0 (Tesla
T4, 14913MiB)
YOLO26n summary (fused): 122 layers, 2,375,616 parameters, 0
gradients, 5.2 GFLOPs
val: Fast image access □ (ping: 0.0±0.0 ms, read: 36.7±3.1 MB/s, size:
58.4 KB)
val: Scanning /content/dataset/valid/labels.cache... 465 images, 5
backgrounds, 0 corrupt: 100% ----- 465/465 130.0Mit/s 0.0s
ages Instances Box(P R mAP50 mAP50-95): 100%
----- 30/30 4.7it/s 6.4s
all 465 1266 0.5 0.614
0.543 0.318
crossing 139 272 0.479 0.675
0.514 0.297
jaywalking 243 331 0.659 0.859

```

0.834	0.518				
	pedestrian	133	396	0.37	0.209
0.251	0.136				
	waiting	152	267	0.493	0.712
0.572	0.32				

Speed: 2.4ms preprocess, 5.2ms inference, 0.0ms loss, 0.3ms postprocess per image
 Results saved to /content/runs/detect/val
 Model 1 cleared from memory

Validating Model 2

Ultralytics 8.4.18 Python-3.12.12 torch-2.10.0+cu128 CUDA:0 (Tesla T4, 14913MiB)
 YOLO26n summary (fused): 122 layers, 2,375,616 parameters, 0 gradients, 5.2 GFLOPs
 val: Fast image access (ping: 0.0±0.0 ms, read: 36.9±22.9 MB/s, size: 66.0 KB)
 val: Scanning /content/dataset/valid/labels.cache... 465 images, 5 backgrounds, 0 corrupt: 100% 465/465 195.0Mit/s 0.0s

ages	Instances	Box(P	R	mAP50	mAP50-95)
30/30	6.0it/s	5.0s			
all	465	1266	0.408	0.539	
0.434	0.254				
	crossing	139	272	0.372	0.607
0.396	0.236				
	jaywalking	243	331	0.606	0.749
0.734	0.468				
	pedestrian	133	396	0.311	0.197
0.206	0.107				
	waiting	152	267	0.343	0.603
0.399	0.207				

Speed: 1.7ms preprocess, 4.0ms inference, 0.0ms loss, 0.2ms postprocess per image
 Results saved to /content/runs/detect/val2
 Model 2 cleared from memory

Validating Model 3

Ultralytics 8.4.18 Python-3.12.12 torch-2.10.0+cu128 CUDA:0 (Tesla T4, 14913MiB)
 YOLO26n summary (fused): 122 layers, 2,375,616 parameters, 0 gradients, 5.2 GFLOPs
 val: Fast image access (ping: 0.0±0.0 ms, read: 133.6±232.2 MB/s, size: 63.1 KB)
 val: Scanning /content/dataset/valid/labels.cache... 465 images, 5 backgrounds, 0 corrupt: 100% 465/465 195.0Mit/s 0.0s

ages	Instances	Box(P	R	mAP50	mAP50-95): 100%
	30/30	4.6it/s 6.6s			
		all	465	1266	0.675 0.66
0.698	0.447				
		crossing	139	272	0.674 0.717
0.735	0.473				
		jaywalking	243	331	0.851 0.865
0.905	0.64				
		pedestrian	133	396	0.515 0.369
0.453	0.256				
		waiting	152	267	0.659 0.689
0.699	0.418				

Speed: 2.5ms preprocess, 4.3ms inference, 0.0ms loss, 0.2ms postprocess per image
 Results saved to /content/runs/detect/val3
 Model 3 cleared from memory

```

# Create overall metrics table
# Initializes an empty list to store the overall metrics for each model
metrics_data = []

# Iterate through each model's validation results
# Collect the overall metrics for each model and store them in the table
for name, data in val_results.items():
    row = {"Model": name} # Start with the model name as the first column
    row.update(data["overall"]) # Add the overall metrics (mAP50, Precision, etc.)
    metrics_data.append(row) # Add the row to the metrics data list

# Create a pandas DataFrame from the metrics data
# The DataFrame will organize the overall performance metrics for easy display
metrics_df = pd.DataFrame(metrics_data)

# Print header for the metrics table
# This adds a visual separator before showing the overall performance metrics
print("=" * 70)
print("OVERALL PERFORMANCE METRICS – ALL 3 MODELS")
print("=" * 70)

# Display the DataFrame containing the overall performance metrics
# This presents the table in a readable format
display(metrics_df)

# Print footer for the metrics table
# This adds a visual separator after displaying the performance

```

```
metrics
print("=" * 70)
```

```
=====
OVERALL PERFORMANCE METRICS – ALL 3 MODELS
=====
```

```
{"summary": "{\n  \"name\": \"metrics_df\",\n  \"rows\": 3,\n  \"fields\": [\n    {\n      \"column\": \"Model\",\n      \"properties\": {\n        \"dtype\": \"string\",\n        \"num_unique_values\": 3,\n        \"samples\": [\n          \"Model 1\",\n          \"Model 2\",\n          \"Model 3\"\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"mAP50\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 0.13302316840811357,\n        \"min\": 0.4335,\n        \"max\": 0.6982,\n        \"num_unique_values\": 3,\n        \"samples\": [\n          0.5427,\n          0.4335,\n          0.6982\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"mAP50-95\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 0.09813444858967722,\n        \"min\": 0.2545,\n        \"max\": 0.447,\n        \"num_unique_values\": 3,\n        \"samples\": [\n          0.3176,\n          0.2545,\n          0.447\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"Precision\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 0.13560259338719646,\n        \"min\": 0.4081,\n        \"max\": 0.6751,\n        \"num_unique_values\": 3,\n        \"samples\": [\n          0.5004,\n          0.4081,\n          0.6751\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"Recall\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 0.06096165461446511,\n        \"min\": 0.539,\n        \"max\": 0.6598,\n        \"num_unique_values\": 3,\n        \"samples\": [\n          0.6137,\n          0.539,\n          0.6598\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"F1 Score\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 0.10180198098923876,\n        \"min\": 0.4645,\n        \"max\": 0.6674,\n        \"num_unique_values\": 3,\n        \"samples\": [\n          0.5513,\n          0.4645,\n          0.6674\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    }\n  ]\n}", "type": "dataframe", "variable_name": "metrics_df"}

# Create per-class metrics table
# Initializes an empty list to store per-class metrics for each model
per_class_rows = []

# Iterate through each model's per-class validation results
```

```

# For each class, store the class-specific metrics in the table
for name, data in val_results.items():
    for cls_data in data["per_class"]:
        row = {"Model": name} # Start with the model name as the
                               # first column
        row.update(cls_data) # Add the per-class metrics (AP50,
                               # Precision, etc.)
        per_class_rows.append(row) # Add the row to the per-class
                                   # metrics data list

# Create a pandas DataFrame from the per-class metrics data
# The DataFrame will organize the per-class performance metrics for
# easy display
per_class_df = pd.DataFrame(per_class_rows)

# Print header for the per-class metrics table
# This adds a visual separator before showing the per-class
# performance metrics
print("=" * 70)
print("    PER-CLASS METRICS – ALL 3 MODELS")
print("=" * 70)

# Display the DataFrame containing the per-class performance metrics
# This presents the table in a readable format
display(per_class_df)
print("=" * 70)

# Create pivot tables for each metric
# For each metric (AP50, Precision, Recall, F1 Score), create a pivot
# table for better visualization
for col in ["AP50", "Precision", "Recall", "F1 Score"]:
    print(f"\n--- {col} by Class ---") # Print the metric header
    # Create a pivot table where 'Class' is the index, 'Model' is the
    # columns, and the metric is the values
    pivot = per_class_df.pivot(
        index="Class", columns="Model", values=col
    )
    # Display the pivot table for the current metric
    display(pivot)

=====
PER-CLASS METRICS – ALL 3 MODELS
=====

{"summary":{"\n  \"name\": \"per_class_df\", \"rows\": 12,\n  \"fields\": [\n    {\n      \"column\": \"Model\", \n      \"properties\": {\n        \"dtype\": \"category\", \n        \"num_unique_values\": 3, \n        \"samples\": [\n          \"Model 1\", \n          \"Model 2\", \n          \"Model 3\" \n        ], \n        \"semantic_type\": \"\", \n        \"description\": \"\" \n      } \n    ] \n  ] \n}}

```



```

n    },\n    {\n        \"column\": \"Class\", \n        \"properties\": {\n            \"dtype\": \"category\", \n            \"num_unique_values\": 4, \n            \"samples\": [\n                \"jaywalking\", \n                \"waiting\", \n                \"crossing\" \n            ], \n            \"semantic_type\": \"\", \n            \"description\": \"\" \n        }, \n        {\n            \"column\": \"AP50\", \n            \"properties\": {\n                \"dtype\": \"number\", \n                \"std\": 0.22615596522325163, \n                \"min\": 0.2056, \n                \"max\": 0.9054, \n                \"num_unique_values\": 12, \n                \"samples\": [\n                    0.4529, \n                    0.9054, \n                    0.5139 \n                ], \n                \"semantic_type\": \"\", \n                \"description\": \"\" \n            }, \n            {\n                \"column\": \"Precision\", \n                \"properties\": {\n                    \"dtype\": \"number\", \n                    \"std\": 0.16541111501255581, \n                    \"min\": 0.3108, \n                    \"max\": 0.8513, \n                    \"num_unique_values\": 12, \n                    \"samples\": [\n                        0.5153, \n                        0.8513, \n                        0.4787 \n                    ], \n                    \"semantic_type\": \"\", \n                    \"description\": \"\" \n                }, \n                {\n                    \"column\": \"Recall\", \n                    \"properties\": {\n                        \"dtype\": \"number\", \n                        \"std\": 0.2271159589017883, \n                        \"min\": 0.197, \n                        \"max\": 0.8648, \n                        \"num_unique_values\": 12, \n                        \"samples\": [\n                            0.3687, \n                            0.8648, \n                            0.6753 \n                        ], \n                        \"semantic_type\": \"\", \n                        \"description\": \"\" \n                    }, \n                    {\n                        \"column\": \"F1 Score\", \n                        \"properties\": {\n                            \"dtype\": \"number\", \n                            \"std\": 0.18960879083155946, \n                            \"min\": 0.2411, \n                            \"max\": 0.858, \n                            \"num_unique_values\": 12, \n                            \"samples\": [\n                                0.4298, \n                                0.858, \n                                0.5603 \n                            ], \n                            \"semantic_type\": \"\", \n                            \"description\": \"\" \n                        } \n                    } \n                } \n            } \n        }, \n        {\n            \"column\": \"AP50\", \n            \"properties\": {\n                \"dtype\": \"number\", \n                \"std\": 0.22615596522325163, \n                \"min\": 0.2056, \n                \"max\": 0.9054, \n                \"num_unique_values\": 12, \n                \"samples\": [\n                    0.4529, \n                    0.9054, \n                    0.5139 \n                ], \n                \"semantic_type\": \"\", \n                \"description\": \"\" \n            }, \n            {\n                \"column\": \"Precision\", \n                \"properties\": {\n                    \"dtype\": \"number\", \n                    \"std\": 0.16541111501255581, \n                    \"min\": 0.3108, \n                    \"max\": 0.8513, \n                    \"num_unique_values\": 12, \n                    \"samples\": [\n                        0.5153, \n                        0.8513, \n                        0.4787 \n                    ], \n                    \"semantic_type\": \"\", \n                    \"description\": \"\" \n                }, \n                {\n                    \"column\": \"Recall\", \n                    \"properties\": {\n                        \"dtype\": \"number\", \n                        \"std\": 0.2271159589017883, \n                        \"min\": 0.197, \n                        \"max\": 0.8648, \n                        \"num_unique_values\": 12, \n                        \"samples\": [\n                            0.3687, \n                            0.8648, \n                            0.6753 \n                        ], \n                        \"semantic_type\": \"\", \n                        \"description\": \"\" \n                    }, \n                    {\n                        \"column\": \"F1 Score\", \n                        \"properties\": {\n                            \"dtype\": \"number\", \n                            \"std\": 0.18960879083155946, \n                            \"min\": 0.2411, \n                            \"max\": 0.858, \n                            \"num_unique_values\": 12, \n                            \"samples\": [\n                                0.4298, \n                                0.858, \n                                0.5603 \n                            ], \n                            \"semantic_type\": \"\", \n                            \"description\": \"\" \n                        } \n                    } \n                } \n            } \n        } \n    ], \n    \"type\": \"dataframe\", \n    \"variable_name\": \"per_class_df\" \n}

```

--- AP50 by Class ---

```

{
  \"summary\": {
    \"name\": \"pivot\",
    \"rows\": 4,
    \"fields\": [
      {
        \"column\": \"Class\",
        \"properties\": {
          \"dtype\": \"string\",
          \"num_unique_values\": 4,
          \"samples\": [
            \"jaywalking\",
            \"waiting\",
            \"crossing\"
          ],
          \"semantic_type\": \"\",
          \"description\": \"\"
        },
        {
          \"column\": \"Model 1\",
          \"properties\": {
            \"dtype\": \"number\",
            \"std\": 0.23931804500845036,
            \"min\": 0.2507,
            \"max\": 0.834,
            \"num_unique_values\": 4,
            \"samples\": [
              0.834,
              0.5722,
              0.5139
            ],
            \"semantic_type\": \"\",
            \"description\": \"\"
          },
          {
            \"column\": \"Model 2\",
            \"properties\": {
              \"dtype\": \"number\",
              \"std\": 0.21969890570202363,
              \"min\": 0.2056,
              \"max\": 0.7339,
              \"num_unique_values\": 4,
              \"samples\": [
                0.7339,
                0.3986,
                0.396
              ],
              \"semantic_type\": \"\",
              \"description\": \"\"
            }
          }
        ]
      }
    ]
  }
}

```



```

\"max\": 0.8587,\n          \"num_unique_values\": 4,\n\"samples\": [\n          0.8587,\n          0.7116,\n          0.6753\n        ],\n          \"semantic_type\": \"\",\n\"description\": \"\"\n        },\n        {\n          \"column\":\n\"Model 2\",\n          \"properties\": {\n          \"dtype\": \"number\",\n          \"std\": 0.23791720548683878,\n          \"min\": 0.197,\n          \"max\": 0.7492,\n          \"num_unique_values\": 4,\n          \"samples\": [\n          0.7492,\n          0.603,\n          0.6066\n        ],\n          \"semantic_type\": \"\",\n          \"description\": \"\"\n        },\n        {\n          \"column\":\n\"Model 3\",\n          \"properties\": {\n          \"dtype\": \"number\",\n          \"std\": 0.208863135011104,\n          \"min\": 0.3687,\n          \"max\": 0.8648,\n          \"num_unique_values\": 4,\n          \"samples\": [\n          0.8648,\n          0.6889,\n          0.7169\n        ],\n          \"semantic_type\": \"\",\n          \"description\": \"\"\n        }\n      ],\n      \"type\": \"dataframe\", \"variable_name\": \"pivot\"}

```

--- F1 Score by Class ---

```

{\"summary\": \"{\\n  \"name\": \"pivot\",\\n  \"rows\": 4,\\n  \"fields\":\n[\\n    {\\n      \"column\": \"Class\",\\n      \"properties\": {\\n\n\"dtype\": \"string\",\\n      \"num_unique_values\": 4,\\n\n\"samples\": [\\n        \"jaywalking\",\\n        \"waiting\",\\n\n\"crossing\",\\n      ],\\n      \"semantic_type\": \"\",\\n\n\"description\": \"\"\n    },\\n    {\\n      \"column\":\n\"Model 1\",\\n      \"properties\": {\\n        \"dtype\": \"number\",\\n\n\"std\": 0.1990853313699095,\\n        \"min\": 0.2674,\\n\n\"max\": 0.7458,\\n        \"num_unique_values\": 4,\\n\n\"samples\": [\\n        0.7458,\\n        0.5828,\\n        0.5603\n      ],\\n        \"semantic_type\": \"\",\\n\n\"description\": \"\"\n    },\\n    {\\n      \"column\":\n\"Model 2\",\\n      \"properties\": {\\n        \"dtype\": \"number\",\\n\n\"std\": 0.17554491115761042,\\n        \"min\": 0.2411,\\n\n\"max\": 0.6703,\\n        \"num_unique_values\": 4,\\n\n\"samples\": [\\n        0.6703,\\n        0.437,\\n        0.4615\n      ],\\n        \"semantic_type\": \"\",\\n\n\"description\": \"\"\n    },\\n    {\\n      \"column\":\n\"Model 3\",\\n      \"properties\": {\\n        \"dtype\": \"number\",\\n\n\"std\": 0.17657522240299364,\\n        \"min\": 0.4298,\\n\n\"max\": 0.858,\\n        \"num_unique_values\": 4,\\n\n\"samples\": [\\n        0.858,\\n        0.6738,\\n        0.6949\n      ],\\n        \"semantic_type\": \"\",\\n\n\"description\": \"\"\n    }\n  ],\n  \"type\": \"dataframe\", \"variable_name\": \"pivot\"}

```

```

# Create bar charts comparing all models across key metrics
# Initializes a subplot with 1 row and 4 columns for displaying the
metrics

```

```

fig, axes = plt.subplots(1, 4, figsize=(22, 5))

# Define the metrics and colors for the bar charts
# metric_cols lists the metrics to be visualized, and colors defines
the bars' colors
metric_cols = ["mAP50", "Precision", "Recall", "F1 Score"]
colors = ["#3498db", "#e74c3c", "#2ecc71"]

# Iterate through each metric and create a bar chart for it
# For each metric, a bar chart is created to visualize model
performance
for idx, col in enumerate(metric_cols):
    # Create a bar chart for the current metric
    # Set the x-axis as the model names and the y-axis as the metric
    values
    bars = axes[idx].bar(
        metrics_df["Model"], metrics_df[col], color=colors
    )

    # Set the title, y-axis limits, and grid for the current chart
    axes[idx].set_title(col, fontsize=14, fontweight="bold")
    axes[idx].set_ylim(0, 1.0) # Limit the y-axis to [0, 1] for
better comparison
    axes[idx].set_ylabel("Score") # Label for the y-axis
    axes[idx].grid(axis="y", alpha=0.3) # Add gridlines on the y-axis
for clarity

    # Add value labels on each bar to show the exact scores
    # This makes the chart more informative by displaying the score at
the top of each bar
    for bar in bars:
        h = bar.get_height()
        axes[idx].annotate(
            f"{h:.4f}", # Display the score with 4 decimal points
            xy=(bar.get_x() + bar.get_width() / 2, h),
            xytext=(0, 5),
            textcoords="offset points",
            ha="center", va="bottom", # Position the label at the top
of the bar
        )

# Set the title for the entire figure
# This provides a general title for the comparison of model
performance
plt.suptitle(
    "Model Performance Comparison",
    fontsize=18, fontweight="bold", y=1.02,
)

# Adjust the layout to ensure the titles and labels fit without

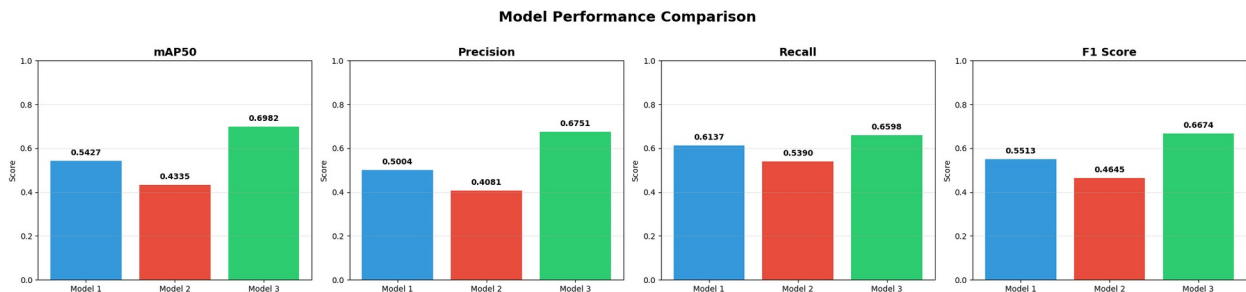
```

```

overlapping
plt.tight_layout()

# Display the final chart with all the bar graphs
plt.show()

```



```

# Create grouped bar chart for per-class F1 scores
# Initializes a figure and axis for the grouped bar chart with a
# specific size
fig, ax = plt.subplots(figsize=(12, 6))

# Set the x positions for the bars based on the number of classes
# Define the width of each bar and the model labels for plotting
x = np.arange(len(class_names))
width = 0.25
model_labels = ["Model 1", "Model 2", "Model 3"]
colors = ["#3498db", "#e74c3c", "#2ecc71"]

# Iterate through each model to plot the corresponding F1 scores for
# each class
# For each model, extract its F1 scores and plot them as bars on the
# grouped chart
for i, mn in enumerate(model_labels):
    vals = per_class_df[
        per_class_df["Model"] == mn
    ]["F1 Score"].values

    # Create bars for each model and class, with appropriate
    # positioning and color
    bars = ax.bar(
        x + i * width, vals, width, label=mn, color=colors[i]
    )

    # Add value labels on each bar to show the F1 score
    # Displays the exact F1 score at the top of each bar for clarity
    for bar in bars:
        h = bar.get_height()
        ax.annotate(
            f"{h:.3f}", # Display the F1 score with 3 decimal places
            xy=(bar.get_x() + bar.get_width() / 2, h),

```

```

        xytext=(0, 3),
        textcoords="offset points",
        ha="center", va="bottom", fontsize=9,
    )

# Set the title for the chart
# Adds a title for the per-class F1 score comparison bar chart
ax.set_title(
    "Per-Class F1 Score Comparison",
    fontsize=16, fontweight="bold",
)

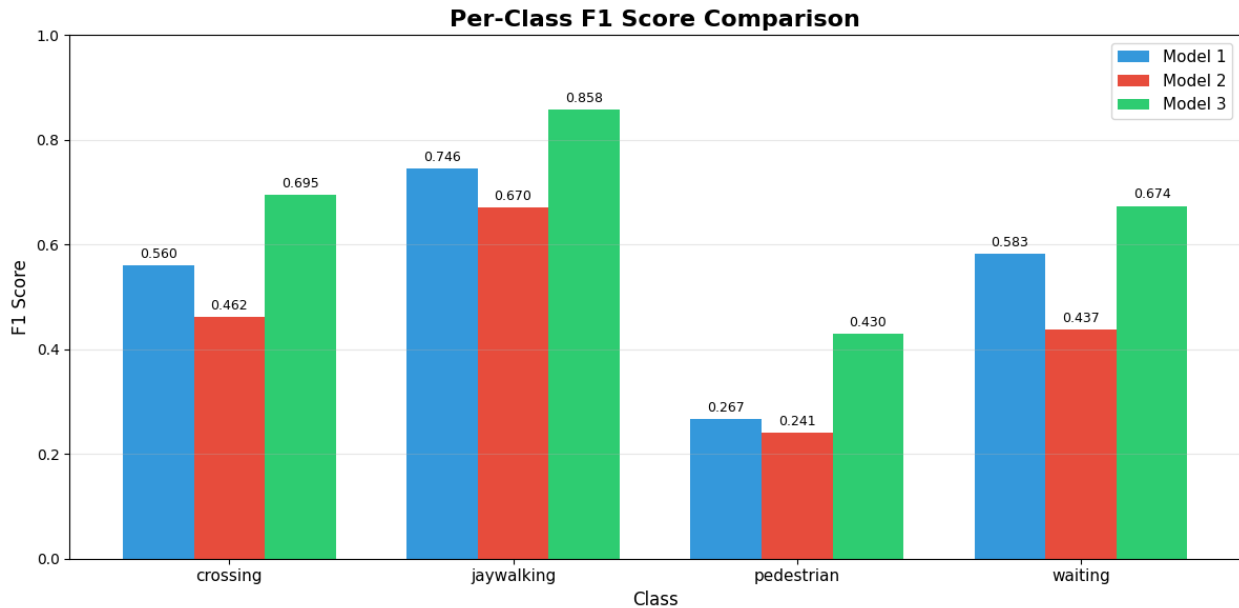
# Set axis labels and tick formatting
# Adds labels for the x-axis, y-axis, and sets the font size for the
# ticks
ax.set_xlabel("Class", fontsize=12)
ax.set_ylabel("F1 Score", fontsize=12)

# Set the positions of the x-ticks and label them with the class names
# This ensures the x-ticks are aligned correctly with the bars
ax.set_xticks(x + width)
ax.set_xticklabels(class_names, fontsize=11)

# Set the y-axis limits and enable gridlines for better readability
# Limits the y-axis from 0 to 1 for normalized F1 scores and adds
# horizontal gridlines
ax.set_ylim(0, 1.0)
ax.legend(fontsize=11)
ax.grid(axis="y", alpha=0.3)

# Adjust the layout to prevent overlap and display the chart
plt.tight_layout()
plt.show()

```



Section 5 - Save the Results

This section outlines how the results from the model evaluation, including performance metrics and hyperparameters, are saved for future use. First, a directory path is defined where the results will be stored, specifically in Google Drive for persistence. The section ensures that the directory exists by creating it if necessary. It then proceeds to save the overall performance metrics and per-class metrics DataFrames as CSV files, which allows for permanent storage and future analysis. Additionally, the hyperparameter settings used for each model are saved in a CSV file, preserving the configuration for reproducibility. This ensures that all essential results and configurations are properly documented and stored in Google Drive for easy access and reference (Google Colab, 2023).

```
# Define the directory path where results will be saved
# This will store all outputs in Google Drive for persistence
save_dir = "/content/drive/MyDrive/DataScience 3/results"

# Create the results directory if it doesn't exist
# exist_ok=True prevents errors if the directory already exists
os.makedirs(save_dir, exist_ok=True)

# Save the overall metrics DataFrame to CSV
# Creates a permanent record of model performance in a CSV file
metrics_df.to_csv(
    os.path.join(save_dir, "overall_metrics.csv"), index=False
)

# Save the per-class metrics DataFrame to CSV
# Preserves detailed class-level performance data for later analysis
per_class_df.to_csv(
    os.path.join(save_dir, "per_class_metrics.csv"), index=False
)
```

```

)

# Save the hyperparameter settings DataFrame to CSV
# Documents the configuration used for each model, useful for
reproducibility
hyperparam_df.to_csv(
    os.path.join(save_dir, "hyperparameter_settings.csv"), index=False
)

# Iterate through each model's training directory
# Copies important files from each model's training folder to the
results folder
for name, dir_path in train_dirs.items():
    # Create a filesystem-safe tag from the model name
    # Converts to lowercase and replaces spaces with underscores for
consistency
    tag = name.lower().replace(" ", "_")

    # Iterate through key visualization files to copy
    # These files contain important training and evaluation plots
    for f in [
        "confusion_matrix.png", # Confusion matrix showing
classification performance
        "confusion_matrix_normalized.png", # Normalized confusion
matrix for clearer comparison
        "results.png", # Visualization of overall training results
    ]:
        print(f"\n Saved to: {save_dir}")

        # Construct the source path for this file
        # Points to the file in the model's training directory
        src = os.path.join(dir_path, f)

        # Check if the source file exists
        # Only copy files that were actually generated during training
        if os.path.exists(src):
            # Save the trained model file for future use
            # Copy the file to the results directory with a model-
specific prefix
            shutil.copy2(
                src, os.path.join(save_dir, f"{tag}_{f}")
            )

            # Print the filename that was saved
            # Provides feedback on what was successfully copied
            print(f" {tag}_{f}")

model_1_confusion_matrix.png
model_1_confusion_matrix_normalized.png
model_1_results.png

```



```
model_1_best.pt
model_2_confusion_matrix.png
model_2_confusion_matrix_normalized.png
model_2_results.png
model_2_best.pt
model_3_confusion_matrix.png
model_3_confusion_matrix_normalized.png
model_3_results.png
model_3_best.pt
```

Saved to: /content/drive/MyDrive/DataScience 3/results

Overall Discussion

The notebook demonstrates the implementation of the YOLOv26 model for object detection using various configurations and hyperparameters, evaluating the performance of different models on a dataset containing images of pedestrians, jaywalking, crossing, and waiting. The models were trained using distinct optimization strategies and hyperparameter settings, including different learning rates, batch sizes, and optimizers (AdamW, SGD, and auto).

The results show varied performance across the three models. Model 1, trained with the AdamW optimizer, exhibited a good overall performance with a mAP50 score of 0.543, and its normalized confusion matrix highlighted strong accuracy in predicting jaywalking and crossing but struggled with pedestrian detection. Model 2, which used the SGD optimizer, showed more consistent performance with a slightly lower mAP50 score of 0.434 but outperformed Model 1 in terms of F1 score for some classes, particularly in detecting jaywalking. Model 3, which utilized automatic optimizer selection, achieved the highest mAP50 score of 0.698, performing robustly across all classes, especially in detecting jaywalking, with an impressive recall and precision for the majority of object classes.

A key aspect of this notebook was the evaluation using confusion matrices, which allowed a detailed comparison of each model's ability to correctly classify the objects in the dataset. The confusion matrices also revealed that the models had difficulty distinguishing between certain classes, notably pedestrians, indicating potential areas for model improvement. Additionally, Model 3's automatic optimizer selection led to superior performance, suggesting that fine-tuning hyperparameters dynamically during training may yield more effective results.

The overall performance metrics, including mAP50, Precision, Recall, and F1 Score, were calculated for each model, with Model 3 outperforming the others in all these key areas. This suggests that the model's ability to adapt its optimizer during training might have contributed to better generalization on unseen data. Furthermore, the hyperparameter table displayed the configurations used in training each model, providing a transparent comparison of the settings across different experiments.

The results indicate the importance of selecting the right optimizer and hyperparameter tuning for maximizing the model's performance in object detection tasks. Model 3's superior performance provides strong evidence that automatic adjustment of hyperparameters can lead to better model training outcomes, especially for complex tasks like object detection in varied environmental conditions.

In conclusion, the notebook effectively demonstrates the utility of YOLOv26 in real-time object detection tasks and the significance of careful hyperparameter optimization. The performance of each model, as evidenced by the training curves, confusion matrices, and evaluation metrics, provides valuable insights into how different configurations affect the model's ability to generalize and classify objects accurately.

References

Google Colab. (2023). Google Colab documentation. <https://colab.research.google.com/>

Pandas Documentation. (2023). Pandas: Python Data Analysis Library.
<https://pandas.pydata.org/pandas-docs/stable/>

Scikit-learn. (2023). Scikit-learn documentation. <https://scikit-learn.org/stable/>

PyTorch Documentation. (2023). PyTorch: Deep Learning Framework.
<https://pytorch.org/docs/stable/>

Matplotlib Documentation. (2023). Matplotlib: Visualization with Python.
<https://matplotlib.org/stable/contents.html>

Ultralytics. (2023). YOLOv5 and YOLOv6 Object Detection with PyTorch.
<https://github.com/ultralytics/yolov5>